

# EE2211 Introduction to Machine Learning

## Lecture 1

Wang Xinchao  
[xinchao@nus.edu.sg](mailto:xinchao@nus.edu.sg)

**Office Hour: Wed 9:30 – 10:30 AM**


# Course Contents

- Introduction and Preliminaries (Xinchao)
  - **Introduction**
  - Data Engineering
  - Introduction to Probability and Statistics
- Fundamental Machine Learning Algorithms I (Vincent)
  - Systems of linear equations
  - Least squares, Linear regression
  - Ridge regression, Polynomial regression
- Fundamental Machine Learning Algorithms II (Vincent)
  - Over-fitting, bias/variance trade-off
  - Optimization, Gradient descent
  - Decision Trees, Random Forest
- Performance and More Algorithms (Xinchao)
  - Performance Issues
  - K-means Clustering
  - Neural Networks

# World's Largest Selfie



# World's Largest Selfie



# Outline

- What is machine learning?
- When do we need machine learning?
- Applications of machine learning
- Types of machine learning
- Walking through a toy example on classification
- Inductive vs. Deductive Reasoning

# What is machine learning?

Learning is any process by which a system improves performance from experience.

- Herbert Simon


A computer program is said to learn

- from **experience E**
- with respect to some class of **tasks T**
- and performance **measure P**,

if its performance at tasks in T, as measured by P, improves with experience E.

- Tom Mitchell

# Machine Learning (Supervised Learning)



**Data      Output**



Cat

:

→  $f(\cdot)$  such that


Dog



$f(\cdot) = \text{'cat'}$

$f(\cdot) = \text{'dog'}$

# Machine Learning (Supervised Learning)



Data      Output



Cat

:



Dog

$$\longrightarrow f(\cdot)$$

When applied



$$f(\text{Cat}) \rightarrow \text{Cat !}$$

New image

Machine Learning: field of study that gives computers the ability to learn without being explicitly programmed

- Arthur Samuel

# AI, Machine Learning, and Deep Learning

## ARTIFICIAL INTELLIGENCE

Any technique which enables computers to mimic human behavior



1950's    1960's    1970's    1980's

## MACHINE LEARNING

AI techniques that give computers the ability to learn without being explicitly programmed to do so



1990's    2000's    2010s

## DEEP LEARNING

A subset of ML which make the computation of multi-layer neural networks feasible



## Example of AI but not ML: Deductive Reasoning

NUS is in Singapore, Singapore is in Asia  $\rightarrow$  NUS is in Asia

# When do we need machine learning?

Lack of human expertise  
 (Navigating on Mars)



Involves huge amount of data  
 (Genomics)



**Learning is not always useful:**

No need to “learn” to calculate payroll!

My Salary = Days\_of\_work \* Daily Salary + Bonus

# Application of Machine Learning

Task T, Performance P, Experience E

T: Digit Recognition


P: Classification Accuracy

E: Labelled Images

4 “four”

3 “three”

Labels -> Supervision!



# Application of Machine Learning

Task T, Performance P, Experience E

T: Email Categorization

P: Classification Accuracy

E: Email Data, Some Labelled




# Application of Machine Learning

Task T, Performance P, Experience E

T: Playing Go Game

P: Chances of Winning

E: Records of Past Games



# Application of Machine Learning


Task T, Performance P, Experience E

T: Identifying Covid-19 Clusters

P: Small Internal Distances

Larger External Distances

E: Records of Patients





*Web Search Engine*



*Product Recommendation*



*Language Translation*



*Photo Tagging*




*Virtual Personal Assistant*



*Portfolio Management*



*Traffic Prediction*



*Medical Diagnosis*



*Algorithmic Trading*

# Types of Machine Learning

## Supervised Learning

Input:

- 1) Training Samples,
- 2) Desired Output  
(Teacher/Supervision)

Output:

A rule that maps input to output

## Unsupervised Learning

Input:

Samples

Output:

Underlying patterns in data

## Reinforcement Learning

Input:

Sequence of States,  
Actions, and  
Delayed Rewards

Output:

Action Strategy: a rule  
that maps the  
environment to action

# Types of Machine Learning

## Supervised Learning

Input:

- 1) Training Samples,
- 2) Desired Output  
(Teacher/Supervision)

Output:

A rule that maps input to output

## Unsupervised Learning

Input:

Samples

Output:

Underlying patterns in data

## Reinforcement Learning


Input:

Sequence of States,  
Actions, and  
Delayed Rewards

Output:



Action Strategy: a rule  
that maps the  
environment to action

# Supervised Learning




## Regression

- Given  $(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_N, y_N)$
- Learn a function  $f(\mathbf{x})$  to predict real-valued  $y$  given  $\mathbf{x}$




# Supervised Learning



## Classification

- Given  $(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_N, y_N)$
- Learn a function  $f(\mathbf{x})$  to predict categorical  $y$  given  $\mathbf{x}$



# Types of Machine Learning

## Supervised Learning

Input:

- 1) Training Samples,
- 2) Desired Output  
(Teacher/Supervision)

Output:

A rule that maps input to output

## Unsupervised Learning

Input:

Samples

Output:

Underlying patterns in data

## Reinforcement Learning

Input:

Sequence of States,  
Actions, and  
Delayed Rewards



Output:

Action Strategy: a rule  
that maps the  
environment to action

# Unsupervised Learning

## Clustering

- Given  $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N$ , without labels
- Output Hidden Structure Behind



# Types of Machine Learning

## Supervised Learning

Input:

- 1) Training Samples,
- 2) Desired Output  
(Teacher/Supervision)

Output:

A rule that maps input to output

## Unsupervised Learning

Input:

Samples

Output:

Underlying patterns in data

## Reinforcement Learning

Input:


Sequence of States,  
Actions, and  
Delayed Rewards

Output:


Action Strategy: a rule  
that maps the  
environment to action

# Reinforcement Learning


## Breakout Game



Initial Performance





Training 15 minutes



Training 30 minutes

# Reinforcement Learning

- Given sequence of states  $S$  and actions  $A$  with (delayed) rewards  $R$
- Output a policy  $\pi(a, s)$ , to guide us what action  $a$  to take in state  $s$



$S$ : Ball Location,  
Paddle Location, Bricks

$A$ : left, right

$R$ :  
positive reward  
 Knocking a brick,  
 clearing all bricks

negative reward  
 Missing the ball

zero reward  
 Cases in between

Supervised  
Unsupervised  
Reinforcement

# Quiz Time!

0 0 0 1 1 1 1 1 2

2 2 2 2 2 2 3 3 3

3 4 4 4 4 5 5 5


6 6 6 7 7 7 7 8 8 8

9 9 9 9 9 9 9 9


Supervised



Supervised




Unsupervised




Reinforcement


# Walking Through A Toy Example: Token Classification



Yes



No



?

?

## Step 1: Feature Extraction

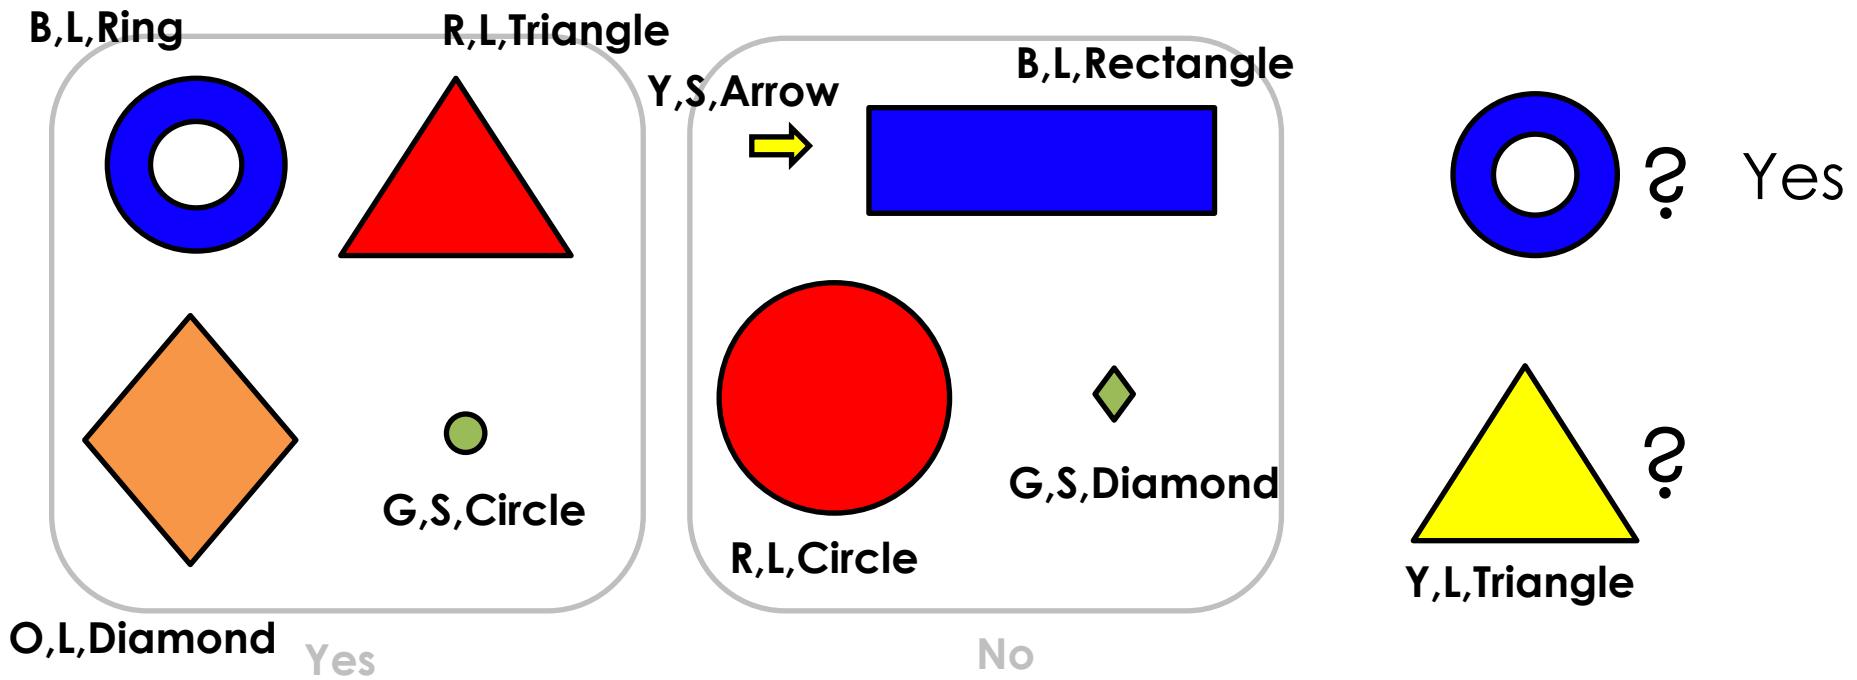
Extract Attributes of Samples



## Step 2: Sample Classification

Decide Label for a Sample

# Walking Through A Toy Example: Token Classification



## Step 1: Feature Extraction

Color, Size, Shape

# Walking Through A Toy Example: Token Classification

## Feature Extraction

	<b>Color</b>	<b>Size</b>	<b>Shape</b>	<b>Label</b>
O	Blue	Large	Ring	Yes
▲	Red	Large	Triangle	Yes
◆	Orange	Large	Diamond	Yes
●	Green	Small	Circle	Yes
→	Yellow	Small	Arrow	No
■	Blue	Large	Rectangle	No
●	Red	Large	Circle	No
◆	Green	Small	Diamond	No
▲	Yellow	Large	Triangle	?

# Walking Through A Toy Example: Token Classification

## Feature Extraction

	Color	Size	Shape	Label
O	Blue	Large	Ring	Yes
▲	Red	Large	Triangle	Yes
◆	Orange	Large	Diamond	Yes
●	Green	Small	Circle	Yes
→	Yellow	Small	Arrow	No
■	Blue	Large	Rectangle	No
○	Red	Large	Circle	No
◆	Green	Small	Diamond	No

# Walking Through A Toy Example: Token Classification

## Feature Extraction

Color	Size	Shape	Label
Blue	Large	Ring	Yes
Red	Large	Triangle	Yes
Orange	Large	Diamond	Yes
Green	Small	Circle	Yes
Yellow	Small	Arrow	No
Blue	Large	Rectangle	No
Red	Large	Circle	No
Green	Small	Diamond	No


## Similarity

Color	Size	Shape	Total
0	1	0	1
0	1	1	2
0	1	0	1
0	0	0	0
1	0	0	1
0	1	0	1
0	1	0	1
0	0	0	0

Legend:

- Yellow triangle: 0
- Red triangle: 1
- Orange diamond: 0
- Green circle: 0
- Yellow arrow: 1
- Blue rectangle: 1
- Red circle: 1
- Green diamond: 0

# Walking Through A Toy Example: Token Classification



**Similarity**


	Color	Size	Shape	Total
▲ ○	0	1	0	1
▲ ▲	0	1	1	2
▲ ◊	0	1	0	1
▲ ●	0	0	0	0
▲ →	1	0	0	1
▲ ■	0	1	0	1
▲ ○	0	1	0	1
▲ ♦	0	0	0	0

## Nearest Neighbor Classifier:


- 1) Find the “nearest neighbor” of a sample in the feature space
- 2) Assign the label of the nearest neighbor to the sample

# Inductive vs. Deductive Reasoning

- Main Task of Machine Learning: to make inference



Probability and Statistics



Rule-based reasoning

NUS is in Singapore, Singapore is in Asia =>  
NUS is in Asia

# Inductive Reasoning

Note: humans use inductive reasoning all the time and not in a formal way like using probability/statistics.

B. C.

by Johnny hart



Ref: Gardner, Martin (March 1979). "[MATHEMATICAL GAMES: On the fabric of inductive logic, and some probability paradoxes](#)" (PDF). *Scientific American*. 234

# Summary by Quick Quiz

## Three Components in ML Definition

Task T, Performance P, Experience E

## Three Types of in ML

Supervised Learning  
Unsupervised Learning  
Reinforcement Learning

## Two Types of Supervised Learning

Classification, Regression

## One Type of Unsupervised Learning

Clustering

## Inductive and Deductive

Inductive: Probable  
Deductive: Rule-based

## Example of a Classifier Model

Nearest Neighbor Classifier

