

EE2211 Lecture 7: Overfitting, Model Complexity, Feature Selection & Regularization, Bias-Variance Tradeoff

Vincent Y. F. Tan



Department of Electrical and Computer Engineering, NUS

EE2211 Spring 2023

Acknowledgements to

Xinchao, Helen, Thomas, Kar Ann, Chen Khong, Robby, and Haizhou

1 Review of Regression

Review of Regression

- Goal: Given feature(s) $\mathbf{x} \in \mathbb{R}^d$, we want to predict target $y \in \mathbb{R}$.
 - 1 \mathbf{x} can be one-dimensional (in which case we may write x);
 - 2 Or \mathbf{x} may be high dimensional;
 - 3 y is often one-dimensional (trying to predict one output).
- Two types of input data
 - 1 **Training Set** $\{(\mathbf{x}_i, y_i)\}_{i=1}^m$;
 - 2 **Test Set** $\{\mathbf{x}_j\}_{j=m+1}^n$ (no targets) [We called one of these samples \mathbf{x}_{new} previously];
- m training samples, n test samples, all in \mathbb{R}^d
- **Learning/Training**
Training set used to estimate regression coefficients $\bar{\mathbf{w}}^*$.
- **Testing/Prediction**
Prediction performed on test set to evaluate performance.

Review of Regression: Linear Case

- Suppose the feature is one-dimensional $x \in \mathbb{R}$.
- Want to learn the affine relationship between x and $y \in \mathbb{R}$.
- With $m = 4$ training samples, the design matrix and target vector are

$$\mathbf{X} = \begin{bmatrix} 1 & x_1 \\ 1 & x_2 \\ 1 & x_3 \\ 1 & x_4 \end{bmatrix} \quad \text{and} \quad \mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ y_3 \\ y_4 \end{bmatrix}$$

■ Learning/Training

$$\bar{\mathbf{w}}^* = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}.$$

■ Testing/Prediction on x_{new}

$$y_{\text{test}} = \begin{bmatrix} 1 \\ x_{\text{new}} \end{bmatrix}^\top \bar{\mathbf{w}}^*.$$

Regression Review: Polynomial

- Features $\mathbf{x} \in \mathbb{R}^d$ may be one-dimensional or high dimensional; $y \in \mathbb{R}$ is a one-dimensional target as usual.
- Want to learn a **polynomial relationship** between x and y
- **Quadratic** illustration ($m = 4$ training samples and x is one-dim) of design matrix and target vector

$$\mathbf{X} = \begin{bmatrix} 1 & x_1 & x_1^2 \\ 1 & x_2 & x_2^2 \\ 1 & x_3 & x_3^2 \\ 1 & x_4 & x_4^2 \end{bmatrix} \quad \text{and} \quad \mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ y_3 \\ y_4 \end{bmatrix}$$

- Learning/Training

$$\bar{\mathbf{w}}^* = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y} \in \mathbb{R}^3.$$

- Testing/Prediction on x_{new}

$$y_{\text{test}} = \begin{bmatrix} 1 \\ x_{\text{new}} \\ x_{\text{new}}^2 \end{bmatrix}^\top \bar{\mathbf{w}}^*.$$

Polynomial Regression Review: Notation

- Sometimes, we use the notation \mathbf{P} in place of \mathbf{X} to emphasize that the features \mathbf{x} are encoded in a polynomial way in the design matrix $\mathbf{P} = \mathbf{X}$.
- Learning/Training

$$\bar{\mathbf{w}}^* = (\mathbf{P}^\top \mathbf{P})^{-1} \mathbf{P}^\top \mathbf{y} \in \mathbb{R}^3.$$

- Testing/Prediction on x_{new}

$$y_{\text{new}} = \underbrace{\begin{bmatrix} 1 \\ x_{\text{new}} \\ x_{\text{new}}^2 \end{bmatrix}}_{=\mathbf{p}_{\text{new}}^\top}^\top \bar{\mathbf{w}}^*.$$

Note on Training & Test Sets

- Affine is special case of polynomial
- Use \mathbf{P} instead of \mathbf{X} from now on.
- Training/Learning (primal) on training set

$$\bar{\mathbf{w}}^* = (\mathbf{P}^\top \mathbf{P})^{-1} \mathbf{P}^\top \mathbf{y} \in \mathbb{R}^{d'}.$$

- Prediction/Testing on the entire test set

$$\mathbf{y}_{\text{new}} = \mathbf{P}_{\text{new}} \bar{\mathbf{w}}^* \in \mathbb{R}^n$$

where

$$\mathbf{P}_{\text{new}} = \begin{bmatrix} \mathbf{p}_{m+1}^\top \\ \mathbf{p}_{m+2}^\top \\ \vdots \\ \mathbf{p}_{m+n}^\top \end{bmatrix} \in \mathbb{R}^{n \times d'}$$


Note on Training & Test Sets

- There should be **zero** overlap between training & test sets.
- Important goal of regression: prediction on **new unseen** data, i.e., the test set.
- Question: Why is test set important for evaluation?

Questions?

Overfitting Motivation: Fukushima Disaster


FIGURE 5-7C: TŌHOKU, JAPAN EARTHQUAKE FREQUENCIES
CHARACTERISTIC FIT



- In March 2011, there was a large earthquake in Japan.
- Small earthquakes occur frequently while massive earthquakes occur rarely.
- Engineers fitted a **polynomial** model
- Earthquake of magnitude 9.0 occurs once every 10,000 years

Overfitting Motivation: Fukushima Disaster

FIGURE 5-7B: TŌHOKU, JAPAN EARTHQUAKE FREQUENCIES
GUTENBERG-RICHTER FIT




Might want to fit an **affine** model instead


Predicts that an earthquake of magnitude 9.0 occurs once every 500 years

Could have reinforced the buildings


Overfitting Example



Overfitting Example




Overfitting Example




	Training Set Fit	
Order 9	Good	

Overfitting Example




	Training Set Fit	Test Set Fit
Order 9	Good	

Overfitting Example



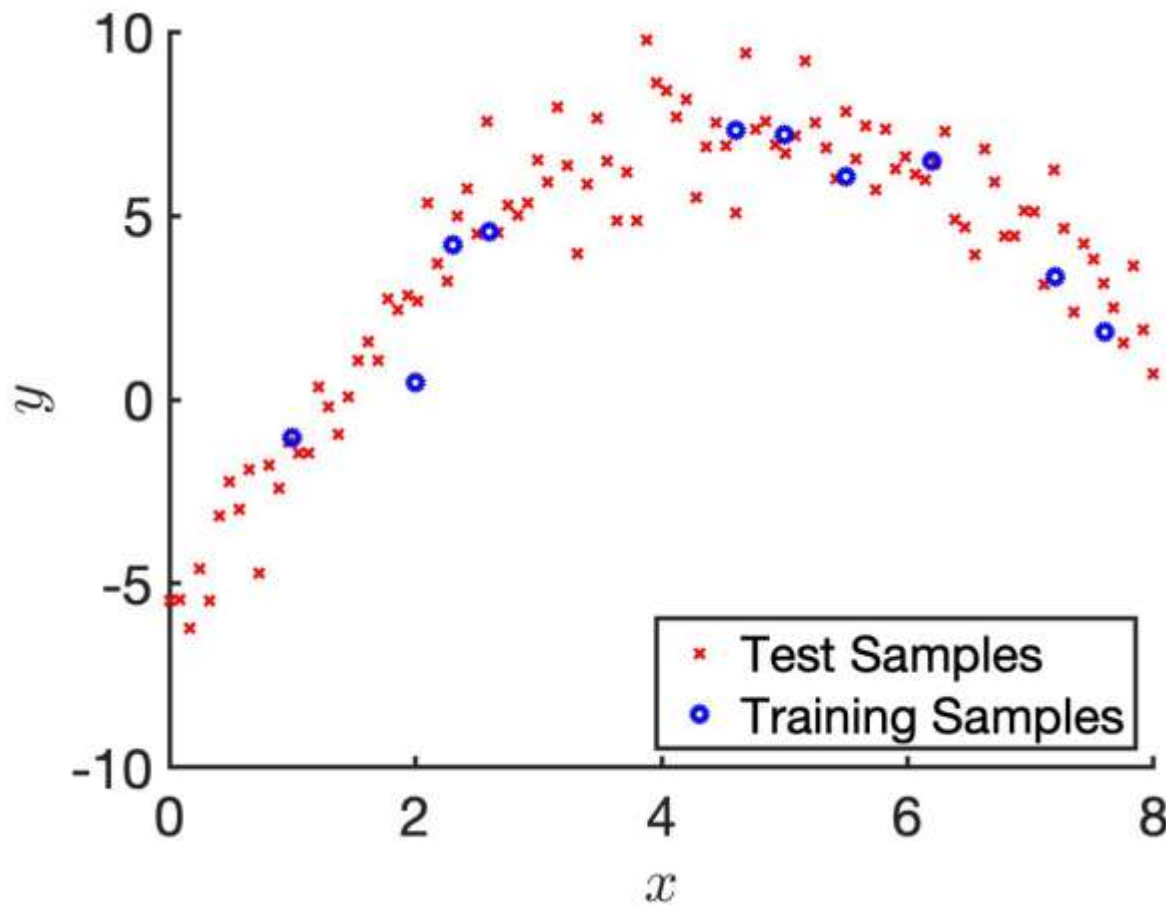
	Training Set Fit	Test Set Fit
Order 9	Good	Bad

Overfitting Example




- If we take one of the blue lines and compute the square of its length, this is called **“squared error”** for that particular data point
- If we average squared errors across all the red crosses, it’s called **mean squared error (MSE)** in the test set

Underfitting Example



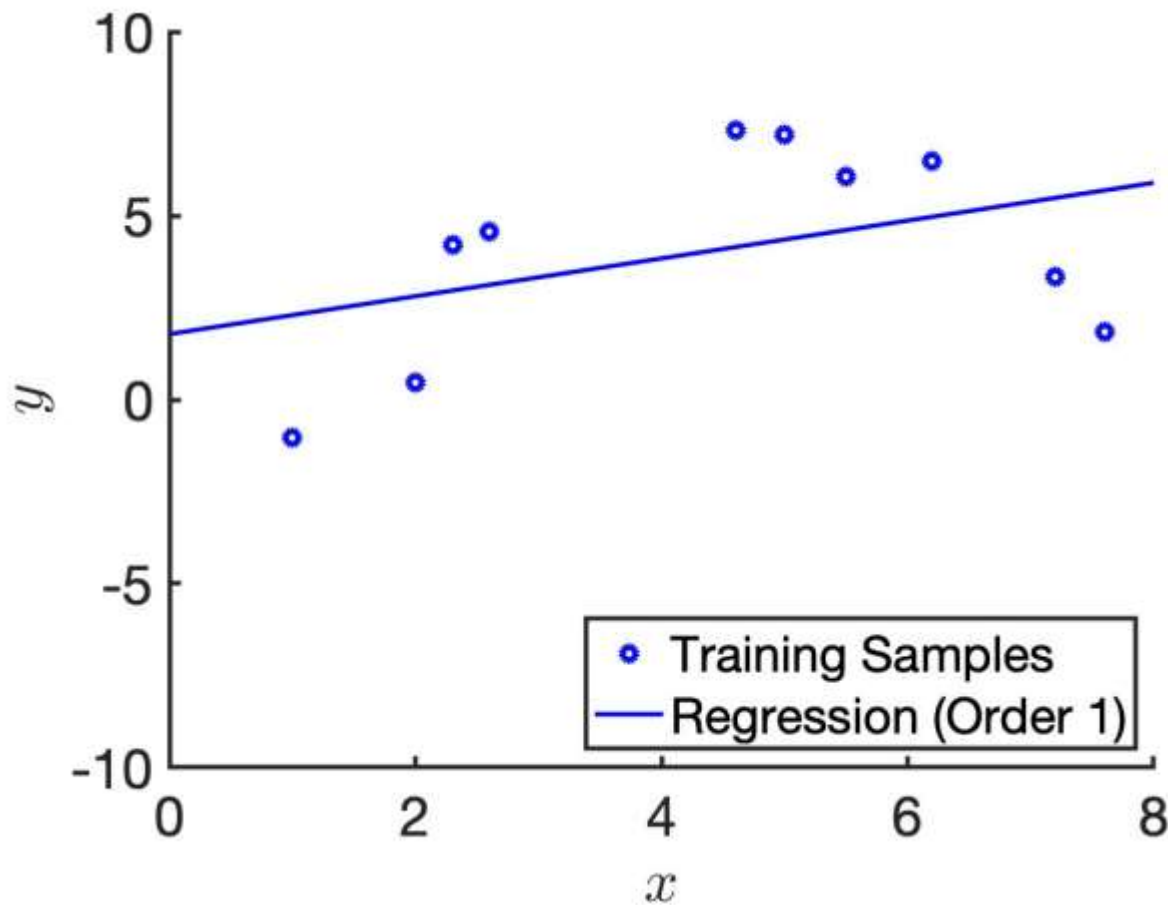
	Training Set Fit	Test Set Fit
Order 9	Good	Bad

Underfitting Example




	Training Set Fit	Test Set Fit
Order 9	Good	Bad

Underfitting Example




	Training Set Fit	Test Set Fit
Order 9	Good	Bad
Order 1		

Underfitting Example




	Training Set Fit	Test Set Fit
Order 9	Good	Bad
Order 1	Bad	

Underfitting Example



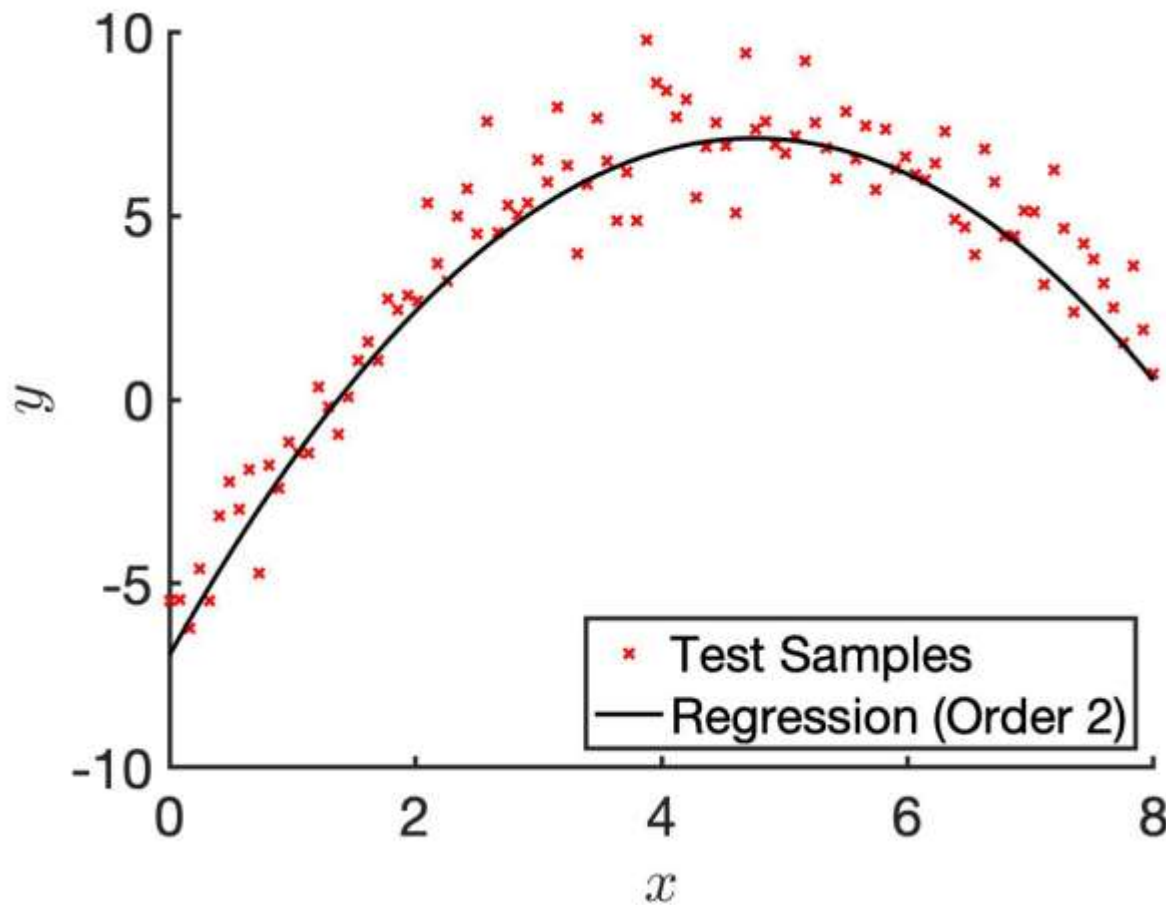
	Training Set Fit	Test Set Fit
Order 9	Good	Bad
Order 1	Bad	Bad

“Just Nice”



	Training Set Fit	Test Set Fit
Order 9	Good	Bad
Order 1	Bad	Bad
Order 2	Good	

“Just Nice”



	Training Set Fit	Test Set Fit
Order 9	Good	Bad
Order 1	Bad	Bad
Order 2	Good	Good

Overfitting & Underfitting

Overfitting



Underfitting



	Training Set Fit	Test Set Fit
Order 9	Good	Bad
Order 1	Bad	Bad
Order 2	Good	Good

Overfitting & Underfitting

- **Overfitting** occurs when model predicts the training data well, but predicts new data (e.g., from test set) poorly

Overfitting & Underfitting

- **Overfitting** occurs when model predicts the training data well, but predicts new data (e.g., from test set) poorly
- **Reason 1**
 - Model is too complex for the data
 - Previous example: Fit order 9 polynomial to 10 data points

Overfitting & Underfitting

- **Overfitting** occurs when model predicts the training data well, but predicts new data (e.g., from test set) poorly
- **Reason 1**
 - Model is too complex for the data
 - Previous example: Fit order 9 polynomial to 10 data points
- **Reason 2**
 - Too many features but number of training samples too small
 - Even linear model can overfit, e.g., linear model with 9 input features (i.e., w is 10-D) and 10 data points in training set => data might not be enough to estimate 10 unknowns well

Overfitting & Underfitting

- **Overfitting** occurs when model predicts the training data well, but predicts new data (e.g., from test set) poorly
- **Reason 1**
 - Model is too complex for the data
 - Previous example: Fit order 9 polynomial to 10 data points
- **Reason 2**
 - Too many features but number of training samples too small
 - Even linear model can overfit, e.g., linear model with 9 input features (i.e., w is 10-D) and 10 data points in training set \Rightarrow data might not be enough to estimate 10 unknowns well
- **Solutions**
 - Use **simpler models** (e.g., lower order polynomial)
 - Use **regularization** (see next part of lecture)

Overfitting & Underfitting

- **Underfitting** is the inability of trained model to predict the targets in the training set


Overfitting & Underfitting

- **Underfitting** is the inability of trained model to predict the targets in the training set
- **Reason 1**
 - Model is too simple for the data
 - Previous example: Fit order 1 polynomial to 10 data points that came from an order 2 polynomial
 - **Solution:** Try more complex model

Overfitting & Underfitting

- **Underfitting** is the inability of trained model to predict the targets in the training set
- **Reason 1**
 - Model is too simple for the data
 - Previous example: Fit order 1 polynomial to 10 data points that came from an order 2 polynomial
 - **Solution:** Try more complex model
- **Reason 2**
 - Features are not informative enough
 - **Solution:** Try to develop more informative features

Overfitting / Underfitting Schematic



Questions?

Feature Selection

- Less features might reduce overfitting
 - Want to discard useless features & keep good features, so perform feature selection

Feature Selection

- Less features might reduce overfitting
 - Want to discard useless features & keep good features, so perform feature selection
- Feature selection procedure
 - Step 1: feature selection in **training** set
 - Step 2: fit model using selected features in **training** set
 - Step 3: evaluate trained model using **test** set

Feature Selection

- Less features might reduce overfitting
 - Want to discard useless features & keep good features, so perform feature selection
- Feature selection procedure
 - Step 1: feature selection in **training** set
 - Step 2: fit model using selected features in **training** set
 - Step 3: evaluate trained model using **test** set
- Very common mistake
 - Feature selection with test set (or full dataset) leads to inflated performance
 - Do not perform feature selection with test data

Selecting Features With Pearson's r

- Given features x , we want to predict target y

Selecting Features With Pearson's r


- Given features x , we want to predict target y
- Assume x & y both continuous

Selecting Features With Pearson's r

- Given features x , we want to predict target y
- Assume x & y both continuous
- Compute Pearson's correlation coefficient between each feature & target y in the training set
 - Pearson's correlation r measures linear relationship between two variables

Selecting Features With Pearson's r


- Given features x , we want to predict target y
- Assume x & y both continuous
- Compute Pearson's correlation coefficient between each feature & target y in the training set
 - Pearson's correlation r measures linear relationship between two variables




Selecting Features With Pearson's r

- Given features x , we want to predict target y
- Assume x & y both continuous
- Compute Pearson's correlation coefficient between each feature & target y in the training set
 - Pearson's correlation r measures linear relationship between two variables


✓ □



✗ □



✓ □



Selecting Features With Pearson's r

- Given features x , we want to predict target y
- Assume x & y both continuous
- Compute Pearson's correlation coefficient between each feature & target y in the training set
 - Pearson's correlation r measures linear relationship between two variables
- Two options
 - Option 1: Pick K features with largest absolute correlations
 - Option 2: Pick all features with absolute correlations $> C$
 - K & C are "magic" numbers set by the ML practitioner

Selecting Features With Pearson's r

- Given features x , we want to predict target y
- Assume x & y both continuous
- Compute Pearson's correlation coefficient between each feature & target y in the training set
 - Pearson's correlation r measures linear relationship between two variables
- Two options
 - Option 1: Pick K features with largest absolute correlations
 - Option 2: Pick all features with absolute correlations $> C$
 - K & C are “magic” numbers set by the ML practitioner
- Other metrics besides Pearson's correlation are possible

Questions?

Regularization

- Regularization is an umbrella term that includes methods forcing learning algorithm to build less complex models.

Regularization

- Regularization is an umbrella term that includes methods forcing learning algorithm to build less complex models.
- **Motivation 1:** Solve an ill-posed problem
 - For example, estimate 10th order polynomial with just 5 datapoints

Regularization

- Regularization is an umbrella term that includes methods forcing learning algorithm to build less complex models.
- **Motivation 1:** Solve an ill-posed problem
 - For example, estimate 10th order polynomial with just 5 datapoints
- **Motivation 2:** Reduce overfitting

Regularization

- Regularization is an umbrella term that includes methods forcing learning algorithm to build less complex models.
- **Motivation 1:** Solve an ill-posed problem
 - For example, estimate 10th order polynomial with just 5 datapoints
- **Motivation 2:** Reduce overfitting
- For example, in previous lecture, we added $\lambda \mathbf{w}^T \mathbf{w}$:

$$\underset{\mathbf{w}}{\operatorname{argmin}} (\mathbf{P}\mathbf{w} - \mathbf{y})^T (\mathbf{P}\mathbf{w} - \mathbf{y}) + \lambda \mathbf{w}^T \mathbf{w}$$

Regularization

- Regularization is an umbrella term that includes methods forcing learning algorithm to build less complex models.
- **Motivation 1:** Solve an ill-posed problem
 - For example, estimate 10th order polynomial with just 5 datapoints
- **Motivation 2:** Reduce overfitting
- For example, in previous lecture, we added $\lambda \mathbf{w}^T \mathbf{w}$:

$$\underset{\mathbf{w}}{\operatorname{argmin}} (\mathbf{P}\mathbf{w} - \mathbf{y})^T (\mathbf{P}\mathbf{w} - \mathbf{y}) + \lambda \mathbf{w}^T \mathbf{w}$$

- Minimizing with respect to \mathbf{w} , primal solution is

$$\hat{\mathbf{w}} = (\mathbf{P}^T \mathbf{P} + \lambda \mathbf{I})^{-1} \mathbf{P}^T \mathbf{y}$$

Regularization

- Regularization is an umbrella term that includes methods forcing learning algorithm to build less complex models.
- **Motivation 1:** Solve an ill-posed problem
 - For example, estimate 10th order polynomial with just 5 datapoints
- **Motivation 2:** Reduce overfitting
- For example, in previous lecture, we added $\lambda \mathbf{w}^T \mathbf{w}$:

$$\underset{\mathbf{w}}{\operatorname{argmin}} (\mathbf{P}\mathbf{w} - \mathbf{y})^T (\mathbf{P}\mathbf{w} - \mathbf{y}) + \lambda \mathbf{w}^T \mathbf{w}$$

- Minimizing with respect to \mathbf{w} , primal solution is

$$\hat{\mathbf{w}} = (\mathbf{P}^T \mathbf{P} + \lambda \mathbf{I})^{-1} \mathbf{P}^T \mathbf{y}$$

- For $\lambda > 0$, matrix becomes invertible (**Motivation 1**)

Regularization

- Regularization is an umbrella term that includes methods forcing learning algorithm to build less complex models.
- **Motivation 1:** Solve an ill-posed problem
 - For example, estimate 10th order polynomial with just 5 datapoints
- **Motivation 2:** Reduce overfitting
- For example, in previous lecture, we added $\lambda \mathbf{w}^T \mathbf{w}$:

$$\underset{\mathbf{w}}{\operatorname{argmin}} (\mathbf{P}\mathbf{w} - \mathbf{y})^T (\mathbf{P}\mathbf{w} - \mathbf{y}) + \lambda \mathbf{w}^T \mathbf{w}$$

- Minimizing with respect to \mathbf{w} , primal solution is

$$\hat{\mathbf{w}} = (\mathbf{P}^T \mathbf{P} + \lambda \mathbf{I})^{-1} \mathbf{P}^T \mathbf{y}$$

- $\hat{\mathbf{w}}$ might also perform better in test set, i.e., reduces overfitting (**Motivation 2**) – will show example later

Regularization

- Consider minimization from previous slide

$$\underset{\mathbf{w}}{\operatorname{argmin}} (\mathbf{P}\mathbf{w} - \mathbf{y})^T (\mathbf{P}\mathbf{w} - \mathbf{y}) + \lambda \mathbf{w}^T \mathbf{w}$$

Regularization

- Consider minimization from previous slide

$$\underset{\mathbf{w}}{\operatorname{argmin}} (\mathbf{P}\mathbf{w} - \mathbf{y})^T (\mathbf{P}\mathbf{w} - \mathbf{y}) + \lambda \mathbf{w}^T \mathbf{w}$$

Cost function quantifying data
fitting error in training set

Regularization

- Consider minimization from previous slide

$$\underset{\mathbf{w}}{\operatorname{argmin}} (\mathbf{P}\mathbf{w} - \mathbf{y})^T (\mathbf{P}\mathbf{w} - \mathbf{y}) + \lambda \mathbf{w}^T \mathbf{w}$$

Cost function quantifying data
fitting error in training set

Regularization

Regularization

- Consider minimization from previous slide

$$\underset{\mathbf{w}}{\operatorname{argmin}} (\mathbf{P}\mathbf{w} - \mathbf{y})^T (\mathbf{P}\mathbf{w} - \mathbf{y}) + \lambda \mathbf{w}^T \mathbf{w}$$

- $\mathbf{w}^T \mathbf{w} = w_0^2 + w_1^2 + \cdots + w_d^2$

Regularization

- Consider minimization from previous slide

$$\underset{\mathbf{w}}{\operatorname{argmin}} (\mathbf{P}\mathbf{w} - \mathbf{y})^T (\mathbf{P}\mathbf{w} - \mathbf{y}) + \lambda \mathbf{w}^T \mathbf{w}$$

- $\mathbf{w}^T \mathbf{w} = w_0^2 + w_1^2 + \cdots + w_d^2$ — L2 - Regularization

Regularization

- Consider minimization from previous slide

$$\underset{\mathbf{w}}{\operatorname{argmin}} (\mathbf{P}\mathbf{w} - \mathbf{y})^T (\mathbf{P}\mathbf{w} - \mathbf{y}) + \lambda \mathbf{w}^T \mathbf{w}$$

- $\mathbf{w}^T \mathbf{w} = w_0^2 + w_1^2 + \cdots + w_d^2$ — L2 - Regularization
- Encourage w_0, \cdots, w_d to be small (also called shrinkage or weight-decay)

Regularization

- Consider minimization from previous slide

$$\underset{\mathbf{w}}{\operatorname{argmin}} (\mathbf{P}\mathbf{w} - \mathbf{y})^T (\mathbf{P}\mathbf{w} - \mathbf{y}) + \lambda \mathbf{w}^T \mathbf{w}$$

- $\mathbf{w}^T \mathbf{w} = w_0^2 + w_1^2 + \cdots + w_d^2$ — L2 - Regularization
- Encourage w_0, \cdots, w_d to be small (also called shrinkage or weight-decay) \Rightarrow constrain model complexity

Regularization

- Consider minimization from previous slide

$$\underset{\mathbf{w}}{\operatorname{argmin}} (\mathbf{P}\mathbf{w} - \mathbf{y})^T (\mathbf{P}\mathbf{w} - \mathbf{y}) + \lambda \mathbf{w}^T \mathbf{w}$$

- $\mathbf{w}^T \mathbf{w} = w_0^2 + w_1^2 + \cdots + w_d^2$ — L2 - Regularization

- Encourage w_0, \dots, w_d to be small (also called shrinkage or weight-decay) \Rightarrow constrain model complexity
- More generally, most machine learning algorithms can be formulated as the following optimization problem

$$\underset{\mathbf{w}}{\operatorname{argmin}} \mathbf{Data-Loss}(\mathbf{w}) + \lambda \mathbf{Regularization}(\mathbf{w})$$

Regularization

- Consider minimization from previous slide

$$\underset{\mathbf{w}}{\operatorname{argmin}} (\mathbf{P}\mathbf{w} - \mathbf{y})^T (\mathbf{P}\mathbf{w} - \mathbf{y}) + \lambda \mathbf{w}^T \mathbf{w}$$

- $\mathbf{w}^T \mathbf{w} = w_0^2 + w_1^2 + \dots + w_d^2$ — L2 - Regularization

- Encourage w_0, \dots, w_d to be small (also called shrinkage or weight-decay) \Rightarrow constrain model complexity
- More generally, most machine learning algorithms can be formulated as the following optimization problem

$$\underset{\mathbf{w}}{\operatorname{argmin}} \mathbf{Data-Loss}(\mathbf{w}) + \lambda \mathbf{Regularization}(\mathbf{w})$$

- Data-Loss(w)** quantifies fitting error to training set given parameters \mathbf{w} : smaller error \Rightarrow better fit to training data

Regularization

- Consider minimization from previous slide

$$\underset{\mathbf{w}}{\operatorname{argmin}} (\mathbf{P}\mathbf{w} - \mathbf{y})^T (\mathbf{P}\mathbf{w} - \mathbf{y}) + \lambda \mathbf{w}^T \mathbf{w}$$


- $\mathbf{w}^T \mathbf{w} = w_0^2 + w_1^2 + \cdots + w_d^2$ — L2 - Regularization

- Encourage w_0, \dots, w_d to be small (also called shrinkage or weight-decay) \Rightarrow constrain model complexity
- More generally, most machine learning algorithms can be formulated as the following optimization problem

$$\underset{\mathbf{w}}{\operatorname{argmin}} \mathbf{Data-Loss}(\mathbf{w}) + \lambda \mathbf{Regularization}(\mathbf{w})$$


- Data-Loss(w)** quantifies fitting error to training set given parameters \mathbf{w} : smaller error \Rightarrow better fit to training data
- Regularization(w)** penalizes more complex models

Regularization Example




	Training Set Fit	Test Set Fit
Order 9	Good	Bad

Regularization Example



	Training Set Fit	Test Set Fit
Order 9	Good	Bad
Order 9, $\lambda = 1$	Good	

Regularization Example




	Training Set Fit	Test Set Fit
Order 9	Good	Bad
Order 9, $\lambda = 1$	Good	Good

Questions?

Bias versus Variance

- Suppose we are trying to predict red target below:




Bias versus Variance

- Suppose we are trying to predict red target below:

Low Bias


High Variance



Low Bias: blue predictions on average close to red target
High Variance: large variability among predictions


Bias versus Variance

- Suppose we are trying to predict red target below:




Bias versus Variance

- Suppose we are trying to predict red target below:



Bias versus Variance

- Suppose we are trying to predict red target below:

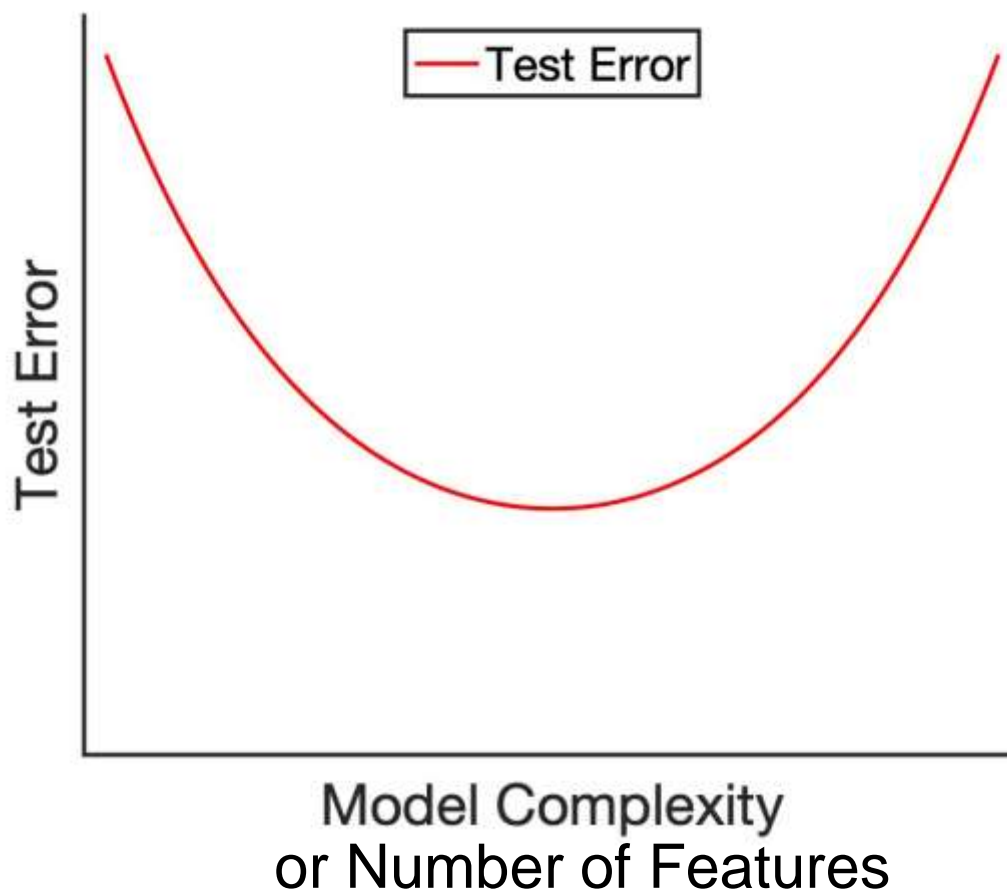


Bias + Variance Trade off

- Test error = Bias Squared + Variance + Irreducible Noise

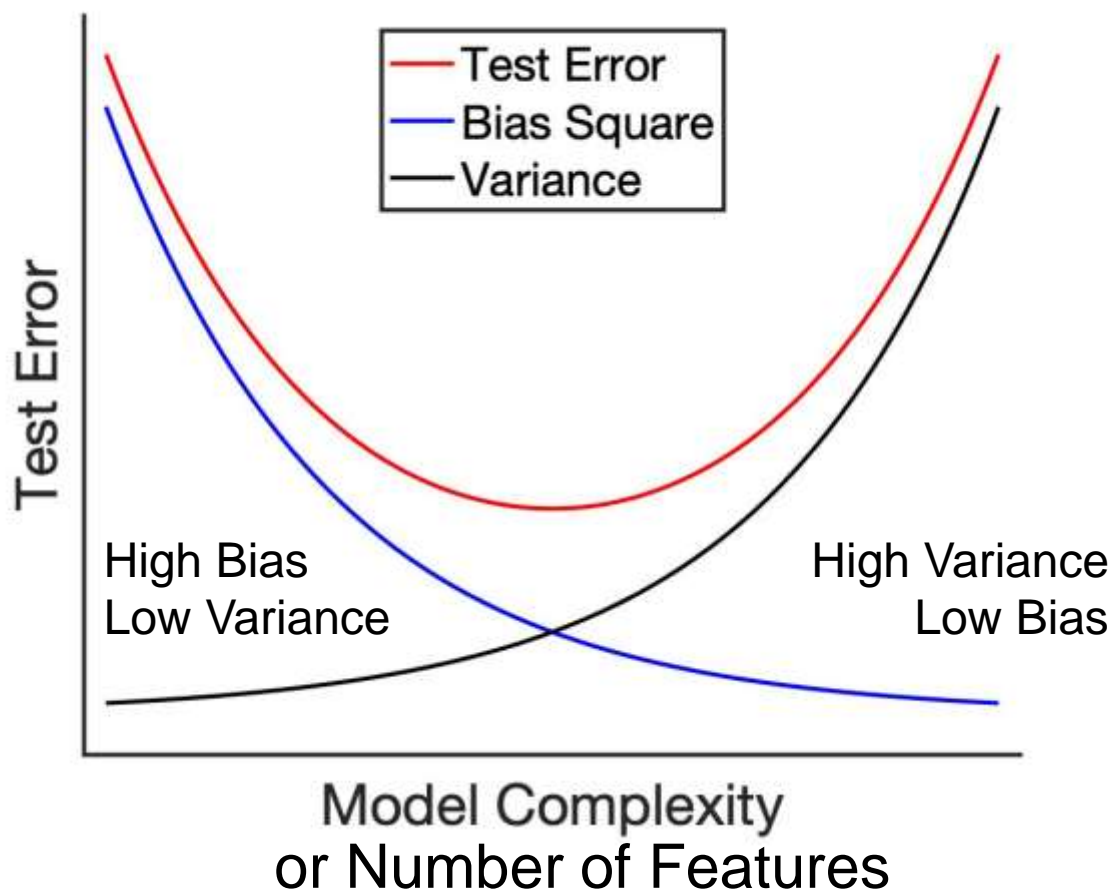
Bias + Variance Trade off

- Test error = Bias Squared + Variance + Irreducible Noise



Bias + Variance Trade off


- Test error = Bias Squared + Variance + Irreducible Noise




Bias + Variance Example

- Simulate data from order 2 polynomial (+ noise)
- Randomly sample 10 training samples each time

4th Order Polynomials




2nd Order Polynomials




Bias + Variance Example

- Simulate data from order 2 polynomial (+ noise)
- Randomly sample 10 training samples each time
- Fit with order 2 polynomial
- Fit with order 4 polynomial

4th Order Polynomials




2nd Order Polynomials




Bias + Variance Example

- Simulate data from order 2 polynomial (+ noise)
- Randomly sample 10 training samples each time
- Fit with order 2 polynomial
- Fit with order 4 polynomial

4th Order Polynomials





2nd Order Polynomials



Bias + Variance Example


- Simulate data from order 2 polynomial (+ noise)
- Randomly sample 10 training samples each time
- Fit with order 2 polynomial
- Fit with order 4 polynomial




Bias + Variance Example

- Simulate data from order 2 polynomial (+ noise)
- Randomly sample 10 training samples each time
- Fit with order 2 polynomial
- Fit with order 4 polynomial

4th Order Polynomials




2nd Order Polynomials




Bias + Variance Example

- Simulate data from order 2 polynomial (+ noise)
- Randomly sample 10 training samples each time
- Fit with order 2 polynomial
- Fit with order 4 polynomial

4th Order Polynomials





2nd Order Polynomials





Bias + Variance Example

- Simulate data from order 2 polynomial (+ noise)
- Randomly sample 10 training samples each time
- Fit with order 2 polynomial: low variance
- Fit with order 4 polynomial: high variance





Bias + Variance Example

- Simulate data from order 2 polynomial (+ noise)
- Randomly sample 10 training samples each time
- Fit with order 2 polynomial: low variance, low bias
- Fit with order 4 polynomial: high variance, low bias



Bias + Variance Example

- Simulate data from order 2 polynomial (+ noise)
 - Randomly sample 10 training samples each time
 - Fit with order 2 polynomial: low variance, low bias
 - Fit with order 4 polynomial: high variance, low bias
- } Order 2
Achieves Lower
Test Error



Bias-Variance Decomposition Theorem

- Test error = Bias Squared + Variance + Irreducible Noise

Bias-Variance Decomposition Theorem

- Test error = Bias Squared + Variance + Irreducible Noise
- Suppose $y(x) = f(x) + \epsilon$, where f is deterministic & ϵ is random with mean 0 & variance σ^2 .

Bias-Variance Decomposition Theorem

- Test error = Bias Squared + Variance + Irreducible Noise
- Suppose $y(x) = f(x) + \epsilon$, where f is deterministic & ϵ is random with mean 0 & variance σ^2 .
- In previous example, we repeated our training 5 times. Each time, we randomly pick 10 different training samples. So training set D is random & there is a probability distribution function $p(D)$ associated with D .

Bias-Variance Decomposition Theorem

- Test error = Bias Squared + Variance + Irreducible Noise
- Suppose $y(x) = f(x) + \epsilon$, where f is deterministic & ϵ is random with mean 0 & variance σ^2 .
- In previous example, we repeated our training 5 times. Each time, we randomly pick 10 different training samples. So training set D is random & there is a probability distribution function $p(D)$ associated with D .
- Using the training data D , we learn a function $\hat{f}_D(x)$. Note dependence of \hat{f}_D on D because \hat{f}_D will change when our training data changes

Bias-Variance Decomposition Theorem

- Test error = Bias Squared + Variance + Irreducible Noise
- Suppose $y(x) = f(x) + \epsilon$, where f is deterministic & ϵ is random with mean 0 & variance σ^2 .
- In previous example, we repeated our training 5 times. Each time, we randomly pick 10 different training samples. So training set D is random & there is a probability distribution function $p(D)$ associated with D .
- Using the training data D , we learn a function $\hat{f}_D(x)$. Note dependence of \hat{f}_D on D because \hat{f}_D will change when our training data changes
- In previous example, we average the predictions across 5 trials. Here, let's perform ∞ trials and average predictions $\hat{f}_D(x)$ across ∞ trials (with different D) resulting in $\hat{f}_{avg}(x) = E_D(\hat{f}_D(x))$

Bias-Variance Decomposition Theorem

- Test error = Bias Squared + Variance + Irreducible Noise
- Suppose $y(x) = f(x) + \epsilon$, where f is deterministic & ϵ is random with mean 0 & variance σ^2 .
- In previous example, we repeated our training 5 times. Each time, we randomly pick 10 different training samples. So training set D is random & there is a probability distribution function $p(D)$ associated with D .
- Using the training data D , we learn a function $\hat{f}_D(x)$. Note dependence of \hat{f}_D on D because \hat{f}_D will change when our training data changes
- In previous example, we average the predictions across 5 trials. Here, let's perform ∞ trials and average predictions $\hat{f}_D(x)$ across ∞ trials (with different D) resulting in $\hat{f}_{avg}(x) = E_D(\hat{f}_D(x))$
- E_D is expectation with respect to $p(D)$: think of this as averaging across ∞ trials

Bias-Variance Decomposition Theorem


- Think of E_D as averaging across ∞ trials

- See proof in extra uploaded material (won't be tested)

Bias-Variance Decomposition Theorem

- Think of E_D as averaging across ∞ trials
- According to the Bias-Variance Decomposition Theorem, the mean squared error for a new test sample x is given by

$$E_D \left[\left(y(x) - \hat{f}_D(x) \right)^2 \right]$$


 \hat{y}_{test}

- See proof in extra uploaded material (won't be tested)

Bias-Variance Decomposition Theorem

- Think of E_D as averaging across ∞ trials
- According to the Bias-Variance Decomposition Theorem, the mean squared error for a new test sample x is given by

$$E_D \left[\left(y(x) - \hat{f}_D(x) \right)^2 \right] = \text{Bias}(\hat{f})^2 + \text{Var}(\hat{f}) + \sigma^2,$$

- See proof in extra uploaded material (won't be tested)

Bias-Variance Decomposition Theorem

- Think of E_D as averaging across ∞ trials
- According to the Bias-Variance Decomposition Theorem, the mean squared error for a new test sample x is given by

$$E_D \left[\left(y(x) - \hat{f}_D(x) \right)^2 \right] = \text{Bias}(\hat{f})^2 + \text{Var}(\hat{f}) + \sigma^2,$$

where

$$\text{Bias}(\hat{f}) = \hat{f}_{avg}(x) - f(x),$$

- See proof in extra uploaded material (won't be tested)

Bias-Variance Decomposition Theorem

- Think of E_D as averaging across ∞ trials
- According to the Bias-Variance Decomposition Theorem, the mean squared error for a new test sample x is given by

$$E_D \left[\left(y(x) - \hat{f}_D(x) \right)^2 \right] = \text{Bias}(\hat{f})^2 + \text{Var}(\hat{f}) + \sigma^2,$$

where

$$\text{Bias}(\hat{f}) = \hat{f}_{avg}(x) - f(x),$$

and

$$\text{Var}(\hat{f}) = E_D \left[\left(\hat{f}_D(x) - \hat{f}_{avg}(x) \right)^2 \right]$$

- See proof in extra uploaded material (won't be tested)

Numerical Example of Bias and Variance Tradeoff I

- We have randomly sampled a training set $D = \{(\mathbf{x}_i, y_i)\}_{i=1}^m$ from a probability distribution $p(D)$.
- Using D , we training a regression model to predict y from \mathbf{x} .
- For simplicity, say that \mathbf{x} and \mathbf{x}_i are one-dimensional, so we write them as x and x_i for brevity.
- We repeat this process of sampling the dataset from $p(D)$ a total of 4 times, resulting in 4 trained models

$$\hat{f}_{D_1}, \hat{f}_{D_2}, \hat{f}_{D_3}, \hat{f}_{D_4}$$

based on the datasets D_1, D_2, D_3, D_4 .

- We have a new test sample $x_{\text{new}} = 3$ whose (unknown) target is $y_{\text{new}} = f(x_{\text{new}}) = 7.5$.
- The function $f(x)$ represents the real/true relation between y and x .

Numerical Example of Bias and Variance Tradeoff II

- Suppose now the predictions of x_{new} from $\hat{f}_{D_1}, \hat{f}_{D_2}, \hat{f}_{D_3}, \hat{f}_{D_4}$ are $\hat{y}_{\text{new},1} = 6$, $\hat{y}_{\text{new},2} = 7$, $\hat{y}_{\text{new},3} = 7$, $\hat{y}_{\text{new},4} = 8$ respectively. That is

$$\hat{y}_{\text{new},1} = \hat{f}_{D_1}(x_{\text{new}}) = \hat{f}_{D_1}(3) = 6$$

$$\hat{y}_{\text{new},2} = \hat{f}_{D_2}(x_{\text{new}}) = \hat{f}_{D_2}(3) = 7$$

$$\hat{y}_{\text{new},3} = \hat{f}_{D_3}(x_{\text{new}}) = \hat{f}_{D_3}(3) = 7$$

$$\hat{y}_{\text{new},4} = \hat{f}_{D_4}(x_{\text{new}}) = \hat{f}_{D_4}(3) = 8$$

- What is the bias and variance based on these models?

Numerical Example of Bias and Variance Tradeoff III

■ Average prediction

$$\hat{f}_{\text{avg}}(x_{\text{new}}) = \hat{f}_{\text{avg}}(3) = \frac{1}{4}(\textcolor{red}{6} + \textcolor{red}{7} + \textcolor{red}{7} + \textcolor{red}{8}) = \textcolor{blue}{7}.$$

■ Bias

$$\text{Bias} = \hat{f}_{\text{avg}}(x_{\text{new}}) - f(x_{\text{new}}) = \hat{f}_{\text{avg}}(x_{\text{new}}) - y_{\text{new}} = \textcolor{blue}{7} - \textcolor{violet}{7.5} = -\frac{1}{2}.$$

■ Variance

$$\begin{aligned}\text{Variance} &= E_D \left[\left(\hat{f}_{\text{avg}}(x_{\text{new}}) - \hat{f}_D(x_{\text{new}}) \right)^2 \right] \\ &= \frac{1}{4} \left((\textcolor{red}{6} - \textcolor{blue}{7})^2 + (\textcolor{red}{7} - \textcolor{blue}{7})^2 + (\textcolor{red}{7} - \textcolor{blue}{7})^2 + (\textcolor{red}{8} - \textcolor{blue}{7})^2 \right) = \frac{1}{2}\end{aligned}$$

Numerical Example of Bias and Variance Tradeoff IV

- If there is no irreducible noise, using the bias-variance formula,

$$\text{MSE} = \text{Bias}^2 + \text{Variance} = \left(-\frac{1}{2}\right)^2 + \frac{1}{2} = \frac{3}{4}.$$

- Now, using actual definition of MSE

$$\begin{aligned}\text{MSE} &= E_D \left[\left(\hat{f}_D(x_{\text{new}}) - f(x_{\text{new}}) \right)^2 \right] \\ &= \frac{1}{4} \left((6 - 7.5)^2 + (7 - 7.5)^2 + (7 - 7.5)^2 + (8 - 7.5)^2 \right) \\ &= \frac{3}{4}\end{aligned}$$

- Voila!

Bias-Variance Python Demo

```
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression, Ridge
from sklearn import datasets

# preparing the dataset into inputs (feature matrix) and outputs (target vector)
data = datasets.load_boston() # fetch the data
X = data.data # feature matrix
y = data.target # target vector

# split the data into training and test samples
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
```

- Here, we are loading the Boston Housing dataset
- We split the raw dataset into 70% training and 30% testing in the last line
- Use .shape to check the dimensions of the various matrices!
- X_train.shape should give (354,13)

Bias-Variance Python Demo

```
def draw_bootstrap_sample(rng, X, y):  
    sample_indices = np.arange(X.shape[0])  
    bootstrap_indices = rng.choice(  
        sample_indices, size=sample_indices.shape[0], replace=True  
    )  
    return X[bootstrap_indices], y[bootstrap_indices]
```

- Here, we are drawing various datasets from the given training set
- We are using a technique known as bootstrapping
- **You don't need to know this!**
- This simulates the effect of ...

Think of E_D as averaging across ∞ trials

Bias-Variance Python Demo

```
def bias_variance_decomp(estimator, X_train, y_train, X_test, y_test, num_rounds=200, random_seed=20):
    rng = np.random.RandomState(random_seed)

    all_pred = []
    for i in range(num_rounds):
        # do bootstrap sampling, i.e., sampling with replacement
        X_boot, y_boot = draw_bootstrap_sample(rng, X_train, y_train)

        # fit a model on bootstrap samples and make prediction on test samples
        pred = estimator.fit(X_boot, y_boot).predict(X_test)
        all_pred.append(pred)

    all_pred = np.array(all_pred)

    # calculate MSE
    avg_mse = ((all_pred - y_test[None,:])**2).mean() # y_test[None,:] will reshape y_test from (N,) to (1,N)

    # average prediction of all bootstrap models on test set
    avg_predictions = np.mean(all_pred, axis=0)

    # calculate bias squared
    avg_bias = np.sum((avg_predictions - y_test)** 2) / y_test.size

    # calculate variance
    avg_var = np.sum((avg_predictions - all_pred)** 2) / all_pred.size

    return avg_mse, avg_bias, avg_var
```

$$E_D \left[\left(y(x) - \hat{f}_D(x) \right)^2 \right]$$

$$\text{Bias}(\hat{f}) = \hat{f}_{avg}(x) - f(x),$$

$$\text{Var}(\hat{f}) = E_D \left[\left(\hat{f}_D(x) - \hat{f}_{avg}(x) \right)^2 \right]$$

- The different samples of the training dataset D are collected in all_pred
- The mean-squared error is calculated in avg_mse
- The bias squared is calculated in avg_bias
- The variance is calculated in avg_var

Bias-Variance Python Demo

```
# define the model
model = LinearRegression()

# estimating the bias and variance
avg_mse, avg_bias, avg_var = bias_variance_decomp(model, X_train,
                                                  y_train, X_test,
                                                  y_test,
                                                  num_rounds=500,
                                                  random_seed=0)

# summary of the results
print('Average mean-squared error: %.3f' % avg_mse)
print('Average bias: %.3f' % avg_bias)
print('Average variance: %.3f' % avg_var)

Average mean-squared error: 27.360
Average bias: 26.064
Average variance: 1.296
```

- Here, we train a linear regression model.
- The MSE is indeed the squared bias plus the variance. Proof by python!
- There is no residual noise in this model.

Bias-Variance Python Demo

```
# define the model
model = Ridge(alpha=1)

# estimating the bias and variance
avg_mse, avg_bias, avg_var = bias_variance_decomp(model, X_train,
                                                  y_train, X_test,
                                                  y_test,
                                                  num_rounds=500,
                                                  random_seed=0)

# summary of the results
print('Average mean-squared error: %.3f' % avg_mse)
print('Average bias: %.3f' % avg_bias)
print('Average variance: %.3f' % avg_var)
model.coef_

Average mean-squared error: 27.375
Average bias: 26.133
Average variance: 1.242
```

- Here, we train a ridge regression model with $\lambda = 1.0$ (it is α in python)
- Bias goes up, variance goes down as expected.
- MSE may go up or down.
- Try experimenting with different α 's (or λ 's)!