

# EE2211 Introduction to Machine Learning

## Lecture 11

Wang Xinchao  
xinchao@nus.edu.sg

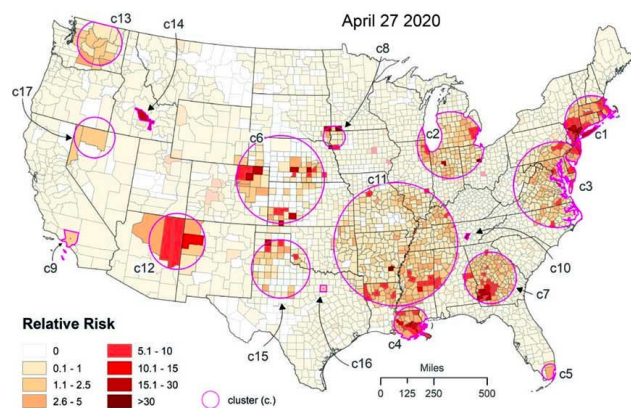
# Course Contents

- Introduction and Preliminaries (Xinchao)
  - Introduction
  - Data Engineering
  - Introduction to Linear Algebra, Probability and Statistics
- Fundamental Machine Learning Algorithms I (Vincent)
  - Systems of linear equations
  - Least squares, Linear regression
  - Ridge regression, Polynomial regression
- Fundamental Machine Learning Algorithms II (Vincent)
  - Over-fitting, bias/variance trade-off
  - Optimization, Gradient descent
  - Decision Trees, Random Forest
- Performance and More Algorithms (Xinchao)
  - Performance Issues
  - **K-means Clustering**
  - Neural Networks

# Outline

- Introduction of unsupervised learning
- K-means Clustering
  - The most popular clustering technique
- Fuzzy Clustering


# Unsupervised Learning




## Discovering Covid clusters



## Business analysis



## Image segmentation



## Community detection in social networks

# Unsupervised Learning

## Introduction

**Motivation:** we do not always have labeled data.

In **unsupervised learning**, the dataset is a collection of **unlabeled examples**  $\{\mathbf{x}_i\}_{i=1}^M$ .

# Unsupervised Learning

## Introduction

Evaluation of unsupervised learning is hard:


- The **absence of labels** representing the desired behavior for your model means **the absence of a solid reference point to judge the quality of your model.**

# Unsupervised Learning

## Main Approaches


- **Clustering**
  - ✓ Groups a set of objects in such a way that objects in the same group (called a **cluster**) are **more similar** (in some sense) to each other than to those in other groups (clusters).
- **Density Estimation**
  - ✓ Models the probability density function (pdf) of the unknown probability distribution from which the dataset has been drawn.
- **Component Analysis**
  - ✓ Breaks down the data from the perspective of signal analysis.
- **Unsupervised Neural Networks**
  - ✓ Autoencoder

# K-means Clustering (2D)

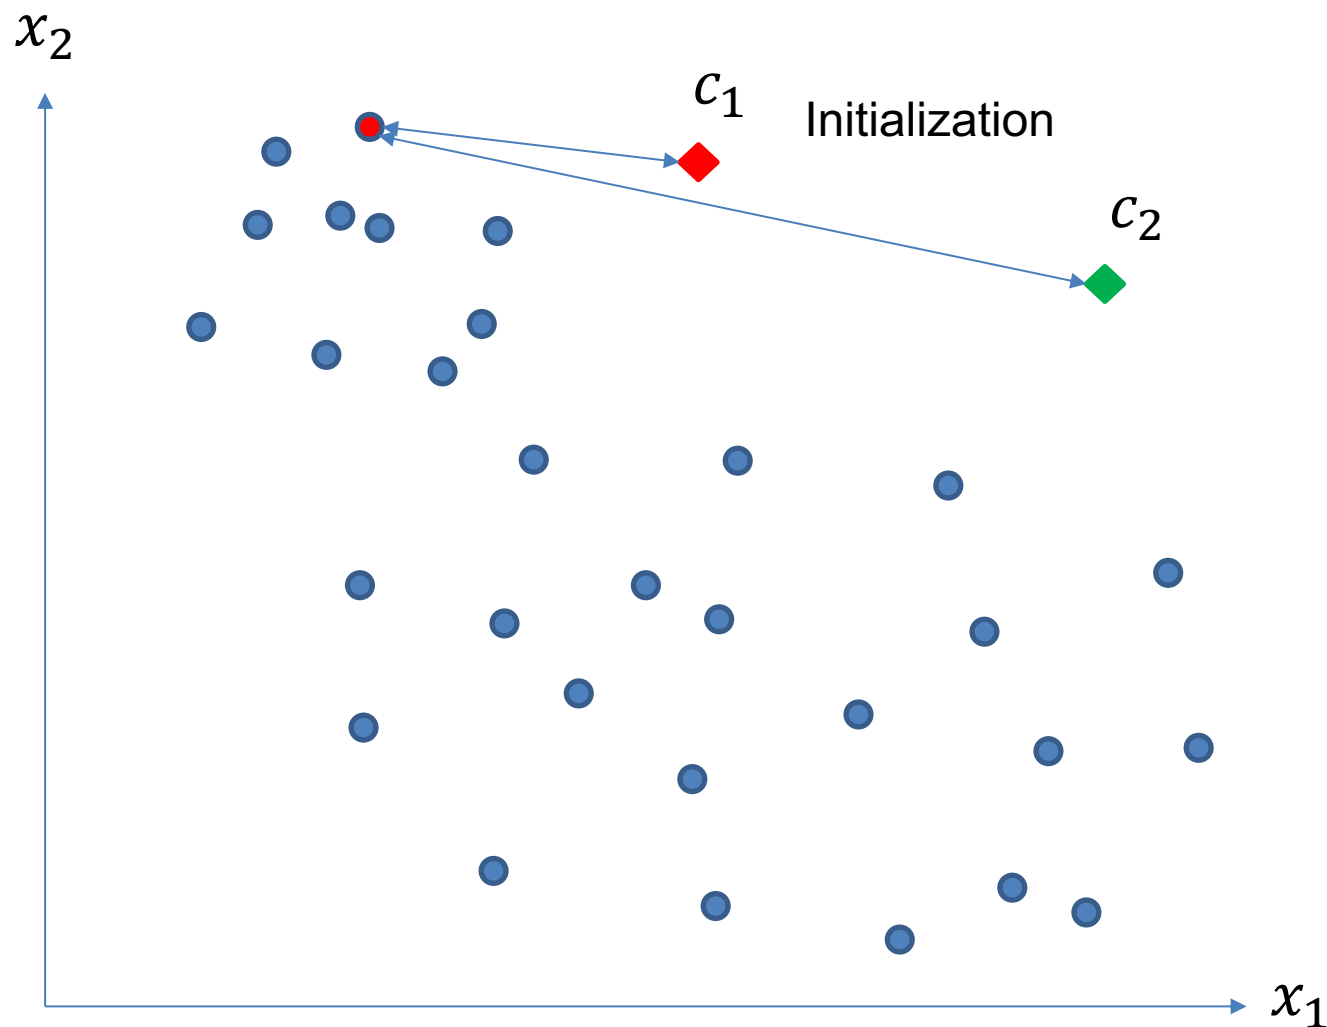




# K-means Clustering




# K-means Clustering




Assignment

# K-means Clustering




Assignment

# K-means Clustering




Assignment

# K-means Clustering




Assignment

# K-means Clustering




Assignment

# K-means Clustering



Assignment


# K-means Clustering



Assignment




# K-means Clustering




Centroid  
Update

# K-means Clustering




# K-means Clustering




Assignment

# K-means Clustering




Assignment

# K-means Clustering




Assignment

# K-means Clustering




Assignment

# K-means Clustering




Centroid  
Update

# K-means Clustering





# K-means Clustering



# K-means Clustering

## Basic/Naïve K-means Clustering

Looping between  
**Assignment** and **Centroid Update**

1. First, we **choose**  $K$  — the number of clusters. Then we randomly select  $K$  feature vectors, called **centroids**, to the feature space.
2. Next, **compute the distance from each example  $x$  to each centroid  $c$**  using some metric, like the Euclidean distance. Then we **assign the closest centroid to each example** (like if we labeled each example with a centroid id as the label).
3. For each centroid, we **calculate the average feature vector** of the examples labeled with it. These average feature vectors become the **new** locations of the **centroids**.
4. We **recompute** the distance from each example to each centroid, modify the assignment and repeat the procedure until **the assignments don't change after the centroid locations are recomputed**.
5. Finally, we **conclude** the clustering with a list of assignments of centroids IDs to the examples.



```
# Define the k-means function
def kmeans_step(data, k, centroids):

    # Assign each data point to the closest centroid
    distances = np.sqrt(((data - centroids[:, np.newaxis])**2).sum(axis=2))
    labels = np.argmin(distances, axis=0)


    # Update centroids to be the mean of the data points assigned to them
    new_centroids = np.zeros_like(centroids)
    for j in range(k):
        new_centroids[j] = np.mean(data[labels == j], axis=0)

    # End if centroids no longer change
    if np.linalg.norm(new_centroids - centroids) < tolerance:
        print("End Clustering, Centroids no change.")
        # Return the original centroids and labels, and set end to True
        return centroids, labels, True
    else:
        # Return the centroids and labels, and set end to False
        return new_centroids, labels, False
```

- See Python code:
  - lec11.ipynb
- See live demo at:
  - lec11\_kmeans.html

**All available in Canvas**  
**Files \ For Students \ Lecture Notes**

# K-means Clustering



[https://en.wikipedia.org/wiki/K-means\\_clustering](https://en.wikipedia.org/wiki/K-means_clustering)

# K-means Clustering

## Optimization Objective Function (within-cluster variance)

$m$ : # of samples;  $i$ : index of samples

$K$ : # of clusters;  $k$ : index of clusters

**Minimize  $J$**

$$J = \sum_{i=1}^m \sum_{k=1}^K w_{ik} \|\mathbf{x}_i - \mathbf{c}_k\|^2 \quad (1)$$

The term  $w_{ik}$  is equal to 1 for data point  $\mathbf{x}_i$  if the data point belongs to cluster  $S_k$ , else  $w_{ik} = 0$ .

Note: The optimization objective function was called  $C(\mathbf{w})$  in Lecture 8. Here, we use  $J$  (with parameters  $w_{ik}$  and  $\mathbf{c}_k$ ) so that it is differentiated from the centroids  $\mathbf{c}_k$ .

Ref: <https://towardsdatascience.com/k-means-clustering-algorithm-applications-evaluation-methods-and-drawbacks-aa03e644b48a>  
[https://en.wikipedia.org/wiki/K-means\\_clustering](https://en.wikipedia.org/wiki/K-means_clustering)

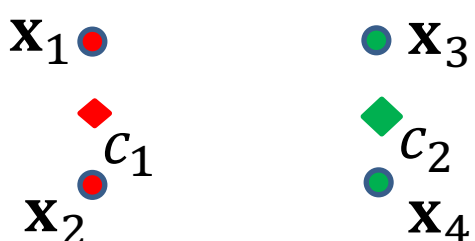
# K-means Clustering

## Optimization Objective Function (within-cluster variance)

Minimize  $J$

$$J = \sum_{i=1}^m \sum_{k=1}^K w_{ik} \|\mathbf{x}_i - \mathbf{c}_k\|^2 \quad (1)$$

$w_{11} = ?$   
 $w_{12} = ?$



$w_{41} = ?$   
 $w_{42} = ?$

The term  $w_{ik}$  is equal to 1 for data point  $\mathbf{x}_i$  if the data point belongs to cluster  $S_k$ , else  $w_{ik} = 0$ .

Note: The optimization objective function was called  $C(\mathbf{w})$  in Lecture 8. Here, we use  $J$  (with parameters  $w_{ik}$  and  $\mathbf{c}_k$ ) so that it is differentiated from the centroids  $\mathbf{c}_k$ .

Ref: <https://towardsdatascience.com/k-means-clustering-algorithm-applications-evaluation-methods-and-drawbacks-aa03e644b48a>  
[https://en.wikipedia.org/wiki/K-means\\_clustering](https://en.wikipedia.org/wiki/K-means_clustering)

# K-means Clustering

## Naïve K-means Algorithm

### 1. Assignment Step (fix $\mathbf{c}$ and update $w$ ):

Computing distances to  
all centroids

$$\mathbf{x}_i \in \mathcal{S}_k \ (w_{ik} = 1) \text{ if } \|\mathbf{x}_i - \mathbf{c}_k\|^2 < \|\mathbf{x}_i - \mathbf{c}_j\|^2 \text{ (else } w_{ik} = 0),$$

$$i = 1, \dots, m; \ j, k = 1, \dots, K.$$

### 2. Update Step (fix $w$ and update $\mathbf{c}$ ):

$$\frac{\partial J}{\partial \mathbf{c}_k} = -2 \sum_{i=1}^m w_{ik} (\mathbf{x}_i - \mathbf{c}_k) = 0 \Rightarrow \mathbf{c}_k = \frac{\sum_{i=1}^m w_{ik} \mathbf{x}_i}{\sum_{i=1}^m w_{ik}}$$

Solving an optimization, i.e., setting derivative to 0

Note:  $\|\mathbf{x} - \mathbf{c}\| = \sqrt{\sum_{d=1}^D (x_d - c_d)^2}$  is called the Euclidean distance.  
where  $\mathbf{x} = (x_1, x_2, \dots, x_D)$ ,  $\mathbf{c} = (c_1, c_2, \dots, c_D)$

# K-means Clustering

## 1. Assignment Step (fix $\mathbf{c}$ and update $w$ ):

$$\mathbf{x}_i \in S_k \text{ } (w_{ik} = 1) \text{ if } \|\mathbf{x}_i - \mathbf{c}_k\|^2 < \|\mathbf{x}_i - \mathbf{c}_j\|^2 \text{ (else } w_{ik} = 0),$$

$$i = 1, \dots, m; \quad j, k = 1, \dots, K.$$

## 2. Update Step (fix $w$ and update $\mathbf{c}$ ):

$$\frac{\partial J}{\partial \mathbf{c}_k} = -2 \sum_{i=1}^m w_{ik} (\mathbf{x}_i - \mathbf{c}_k) = 0 \Rightarrow \mathbf{c}_k = \frac{\sum_{i=1}^m w_{ik} \mathbf{x}_i}{\sum_{i=1}^m w_{ik}}$$


By repeating this two steps, the total loss  $J = \sum_{i=1}^m \sum_{k=1}^K w_{ik} \|\mathbf{x}_i - \mathbf{c}_k\|^2$ , is **guaranteed to NOT increase (i.e., remain the same or decrease)** until convergence.

Why? **At Step 2:** we compute the new mean, by solving an optimization, i.e., compute the derivative and set to zero, and solve  $\mathbf{c}_k$ . This means that, the new  $\mathbf{c}_k$  is guaranteed to give a smaller  $J$  value.


**At Step 1:** we only change the assignment, if the distance to the new centroid is smaller! In other words, we either remain in the old group, or change to a new group that is closer (i.e., gives a smaller  $J$ )




# K-means Clustering (1 D)



# K-means Clustering (1 D)




# K-means Clustering (1 D)




**Different initializations give different clusters!**

# K-means Clustering

- Unfortunately, k-means is not guaranteed to find a global minimum, it finds only local minimum.
- Example:



K-means



Optimal  $J$

- Finding the optimal  $J$  is NP-hard\*
- In practice, k-means clustering usually performs well
- It can be very efficient, and its solution can be used as a starting point for other clustering algorithms

\*<https://en.wikipedia.org/wiki/NP-hardness>

# K-means Clustering


## • Initialization

Initialization by centroid

### Forgy method:

- Randomly chooses  $k$  observations from the dataset and uses these as the initial means.

$k=3$




### Random partition:

- First randomly assigns a cluster to each observation and then proceeds to the update step, thus computing the initial mean to be the centroid of the cluster's randomly assigned points

Initialization by grouping

$k=3$



Ref: [https://en.wikipedia.org/wiki/K-means\\_clustering#Standard\\_algorithm\\_\(naive\\_k-means\)](https://en.wikipedia.org/wiki/K-means_clustering#Standard_algorithm_(naive_k-means))

# Hard vs Soft Clustering

## Hard clustering:

Each data point can belong only one cluster, e.g. K-means

- For example, an apple can be red **OR** green (hard clustering)

## Soft clustering (also known as Fuzzy clustering):

Each data point can belong to more than one cluster.

- For example, an apple can be red **AND** green (fuzzy clustering)
- Here, the apple can be red to a certain degree as well as green to a certain degree.
- Instead of the apple belonging to green [green = 1] and not red [red = 0], the apple can belong to green [green = 0.5] and red [red = 0.5]. These value are normalized between 0 and 1; however, they do not represent probabilities, so the two values **do not need to add up to 1**.

Ref: [https://en.wikipedia.org/wiki/Fuzzy\\_clustering](https://en.wikipedia.org/wiki/Fuzzy_clustering)

# Hard vs Soft Clustering

## Objective Function for Fuzzy C-means

Minimize  $J$

$$J = \sum_{i=1}^m \sum_{k=1}^C (w_{ik})^r \|\mathbf{x}_i - \mathbf{c}_k\|^2$$

$$\text{where } w_{ik} = \frac{1}{\sum_{j=1}^C \left( \frac{\|\mathbf{x}_i - \mathbf{c}_k\|}{\|\mathbf{x}_i - \mathbf{c}_j\|} \right)^{\frac{2}{r-1}}}$$

**No need to memorize**

Each element,  $w_{ik} \in [0,1]$ , tells the degree to which element,  $\mathbf{x}_i$ , belongs to cluster  $\mathbf{c}_k$ .

The fuzzifier  $r > 1$  determines the level of cluster fuzziness; usually  $1.25 \leq r \leq 2$ .

# Hard vs Soft Clustering

## Objective Function for Fuzzy C-means


Minimize  $J$

$$J = \sum_{i=1}^m \sum_{k=1}^C (w_{ik})^r \|\mathbf{x}_i - \mathbf{c}_k\|^2$$

$$\text{where } w_{ik} = \frac{1}{\sum_{j=1}^C \left( \frac{\|\mathbf{x}_i - \mathbf{c}_k\|}{\|\mathbf{x}_i - \mathbf{c}_j\|} \right)^{\frac{2}{r-1}}}$$

$$w_{11} = 0.6$$

$$w_{12} = 0.2$$



$$w_{41} = 0.18$$


$$w_{42} = 0.75$$

Each element,  $w_{ik} \in [0,1]$ , tells the degree to which element,  $\mathbf{x}_i$ , belongs to cluster  $\mathbf{c}_k$ .


The fuzzifier  $r > 1$  determines the level of cluster fuzziness; usually  $1.25 \leq r \leq 2$ .



Iteration 1




Iteration 2




## Visualization of Fuzzy C-means Iterations


Iteration 3



Iteration 4




Iteration 5




## Visualization of Fuzzy C-means Iterations


Iteration 6



Iteration 7



Iteration 8





## Visualization of Fuzzy C-means Iterations

# Hard vs Soft Clustering

## Naïve K-means versus Fuzzy C-means

Naïve K-means:  $w_{ik} \in \{0,1\}$

Fuzzy C-means:  $w_{ik} \in [0,1]$



Ref: [https://en.wikipedia.org/wiki/Fuzzy\\_clustering](https://en.wikipedia.org/wiki/Fuzzy_clustering)

# Summary

- Introduction of unsupervised learning
- K-means Clustering
  - The most popular clustering technique
- Fuzzy Clustering

