

EE2211 Introduction to Machine Learning

Lecture 12

Wang Xinchao
xinchao@nus.edu.sg

Course Contents

- Introduction and Preliminaries (Xinchao)
 - Introduction
 - Data Engineering
 - Introduction to Linear Algebra, Probability and Statistics
- Fundamental Machine Learning Algorithms I (Vincent)
 - Systems of linear equations
 - Least squares, Linear regression
 - Ridge regression, Polynomial regression
- Fundamental Machine Learning Algorithms II (Vincent)
 - Over-fitting, bias/variance trade-off
 - Optimization, Gradient descent
 - Decision Trees, Random Forest
- Performance and More Algorithms (Xinchao)
 - Performance Issues
 - K-means Clustering
 - Neural Networks

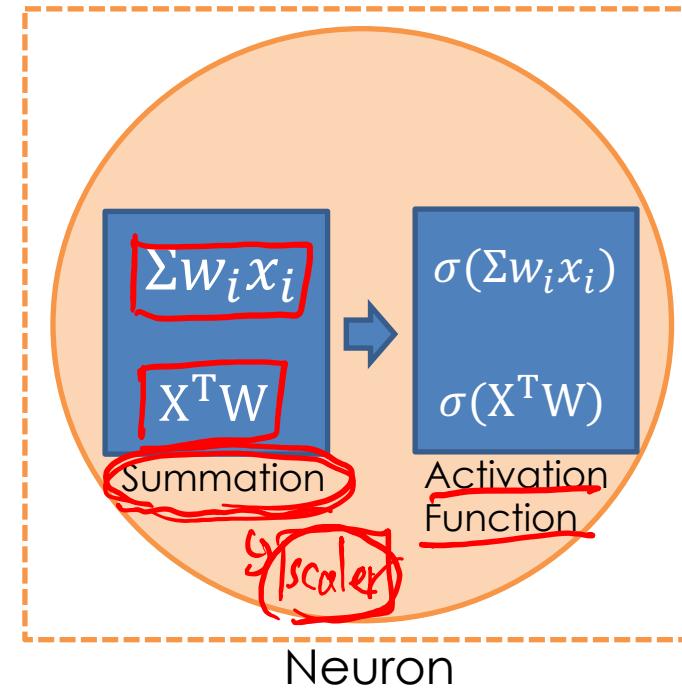
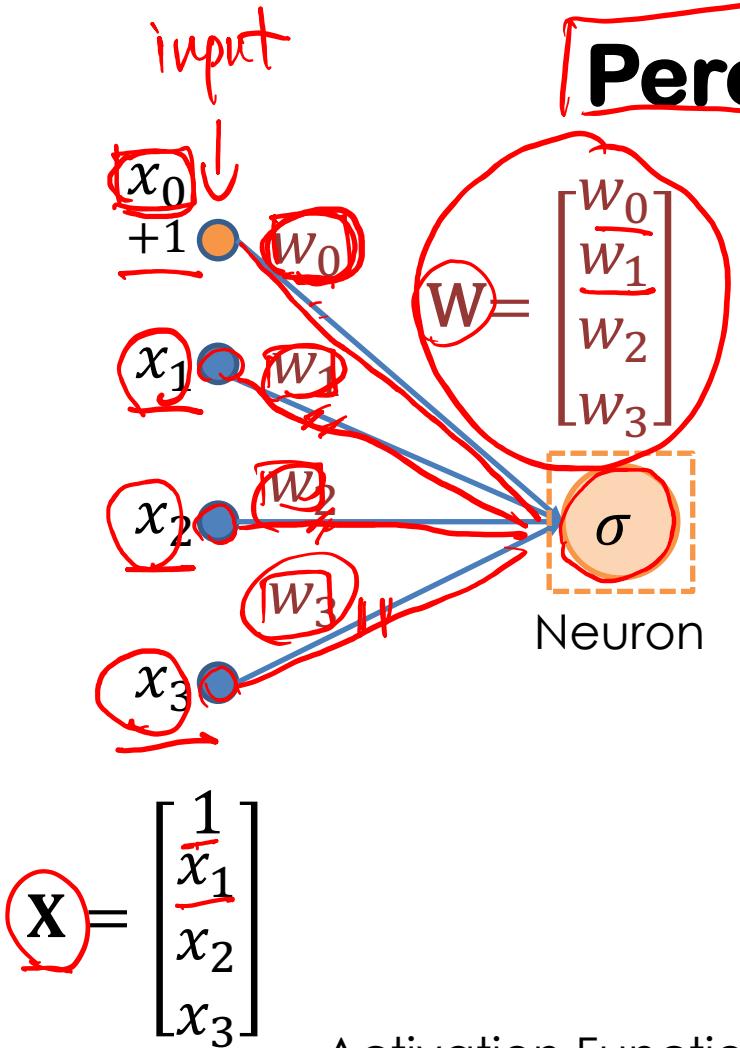
About this week's lecture...

- Neural Network (NN) is a very big topic
 - In NUS we have multiple full-semester modules to discuss NN
 - EE4305 Fuzzy/Neural Systems for Intelligent Robotics
 - EE5934/EE6934 Deep Learning
 - ...
 - In EE2211, we only give a very gentle introduction
- Understanding at conceptual level is sufficient
 - In final exam, we have only 1 True/False + 1 MCQ about NN
 - No computation is required
- You will do some computation in tutorial, but final exam will be much simpler than the questions in tutorial

Outline

- Introduction to Neural Networks
 - Perceptron
 - Activation Functions
 - Multi-layer Perceptron
- Training and Testing of Neural Networks
 - Training: Forward and Backward
 - Testing: Forward
- Convolutional Neural Networks (CNNs)

Perceptron



Output of Neuron: $\sigma(X^T W)$ or $\sigma(\sum w_i x_i)$

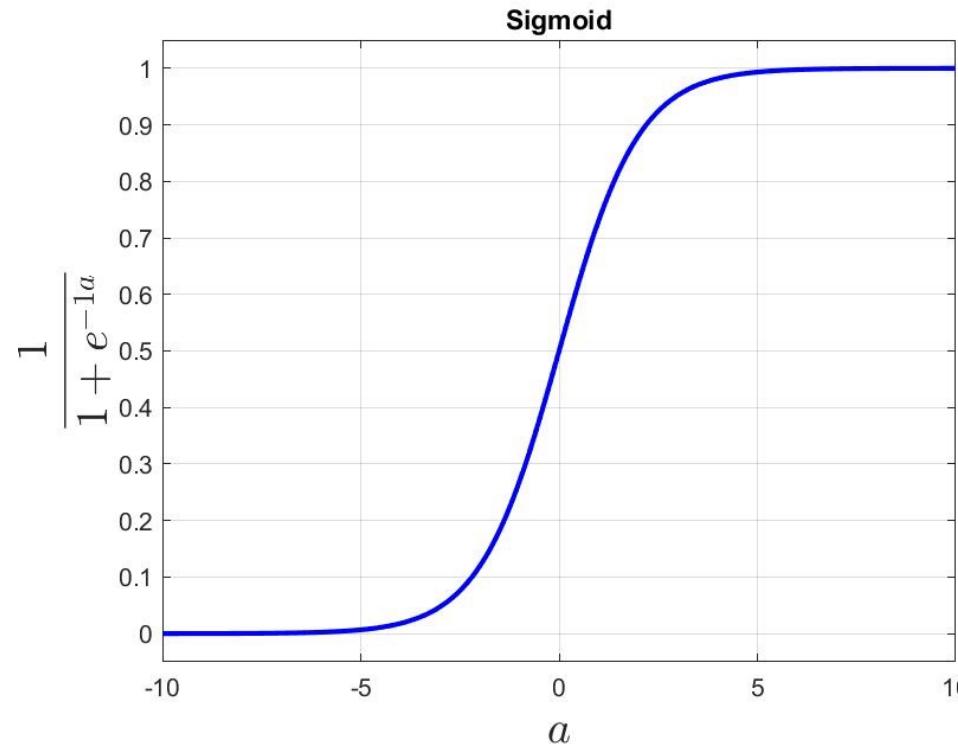
Activation Function: non-linear function to introduce non-linearity into the neural networks!

Goal of training: to learn W !

Activation Functions

Sigmoid Activation Function

$$\sigma(a) = \frac{1}{1 + e^{-\beta a}},$$

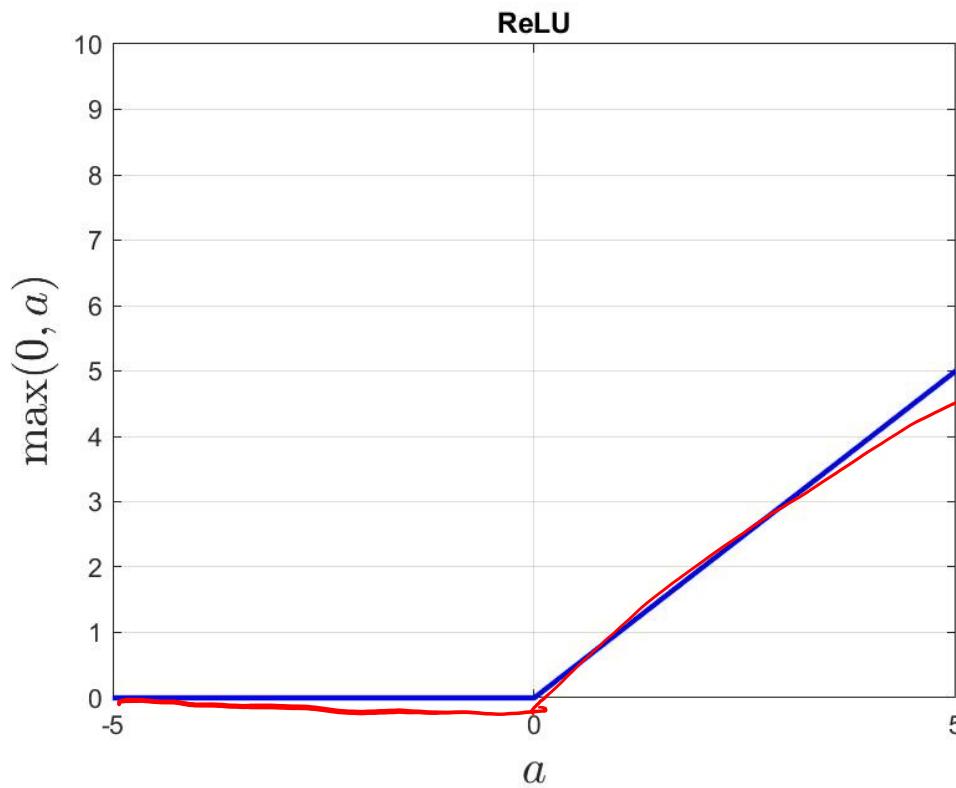


Activation Functions

ReLU Activation Function

$$\sigma(a) = \max(0, a)$$

Rectified Linear Unit (ReLU)



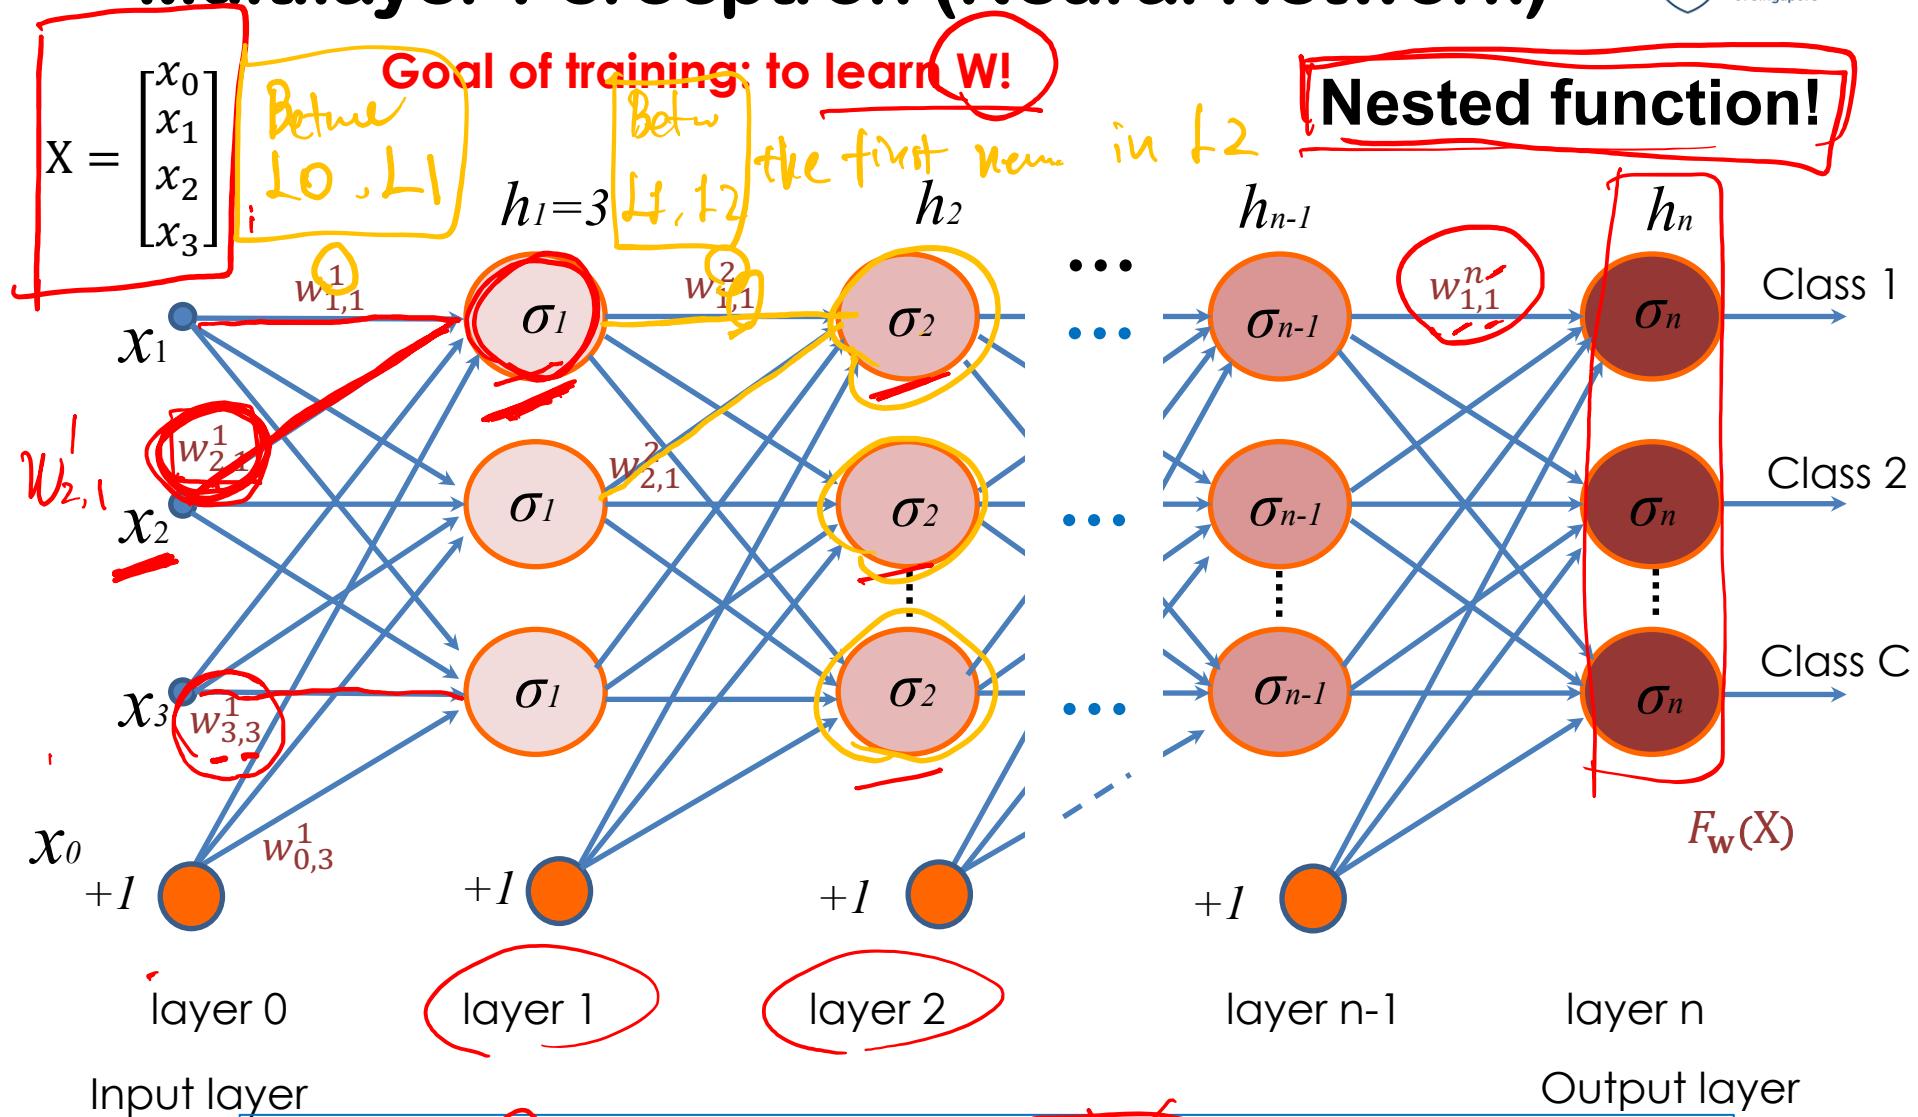
$$a = 1.5$$

$$\sigma(a) = 1.5$$

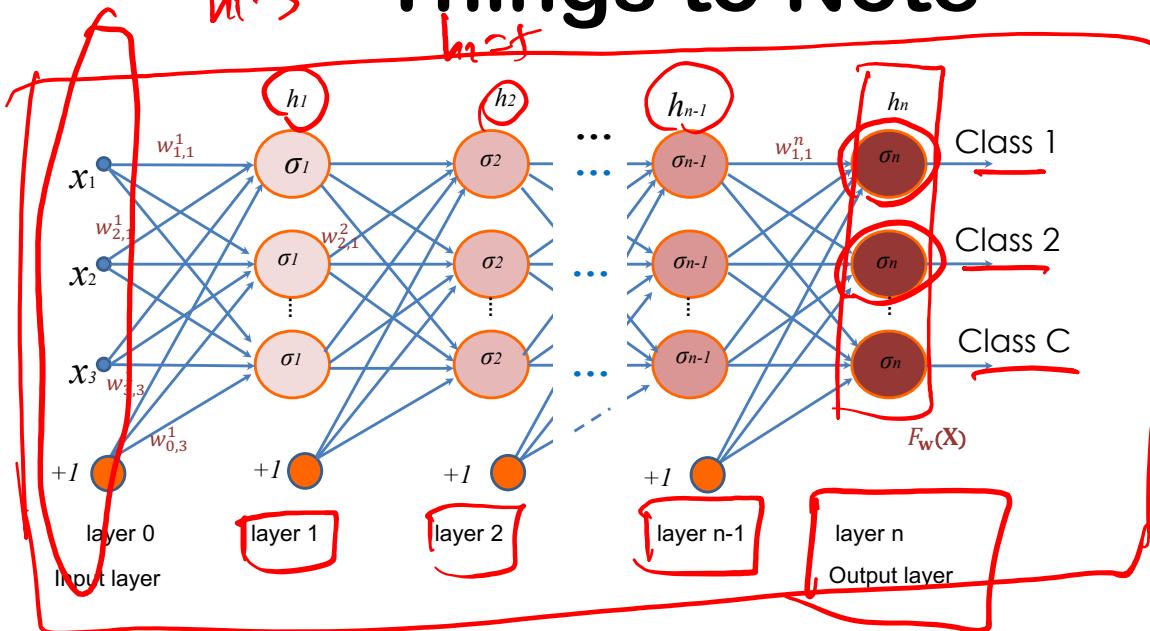
$$a = -2$$

$$\sigma(a) = 0$$

Multilayer Perceptron (Neural Network)



Things to Note

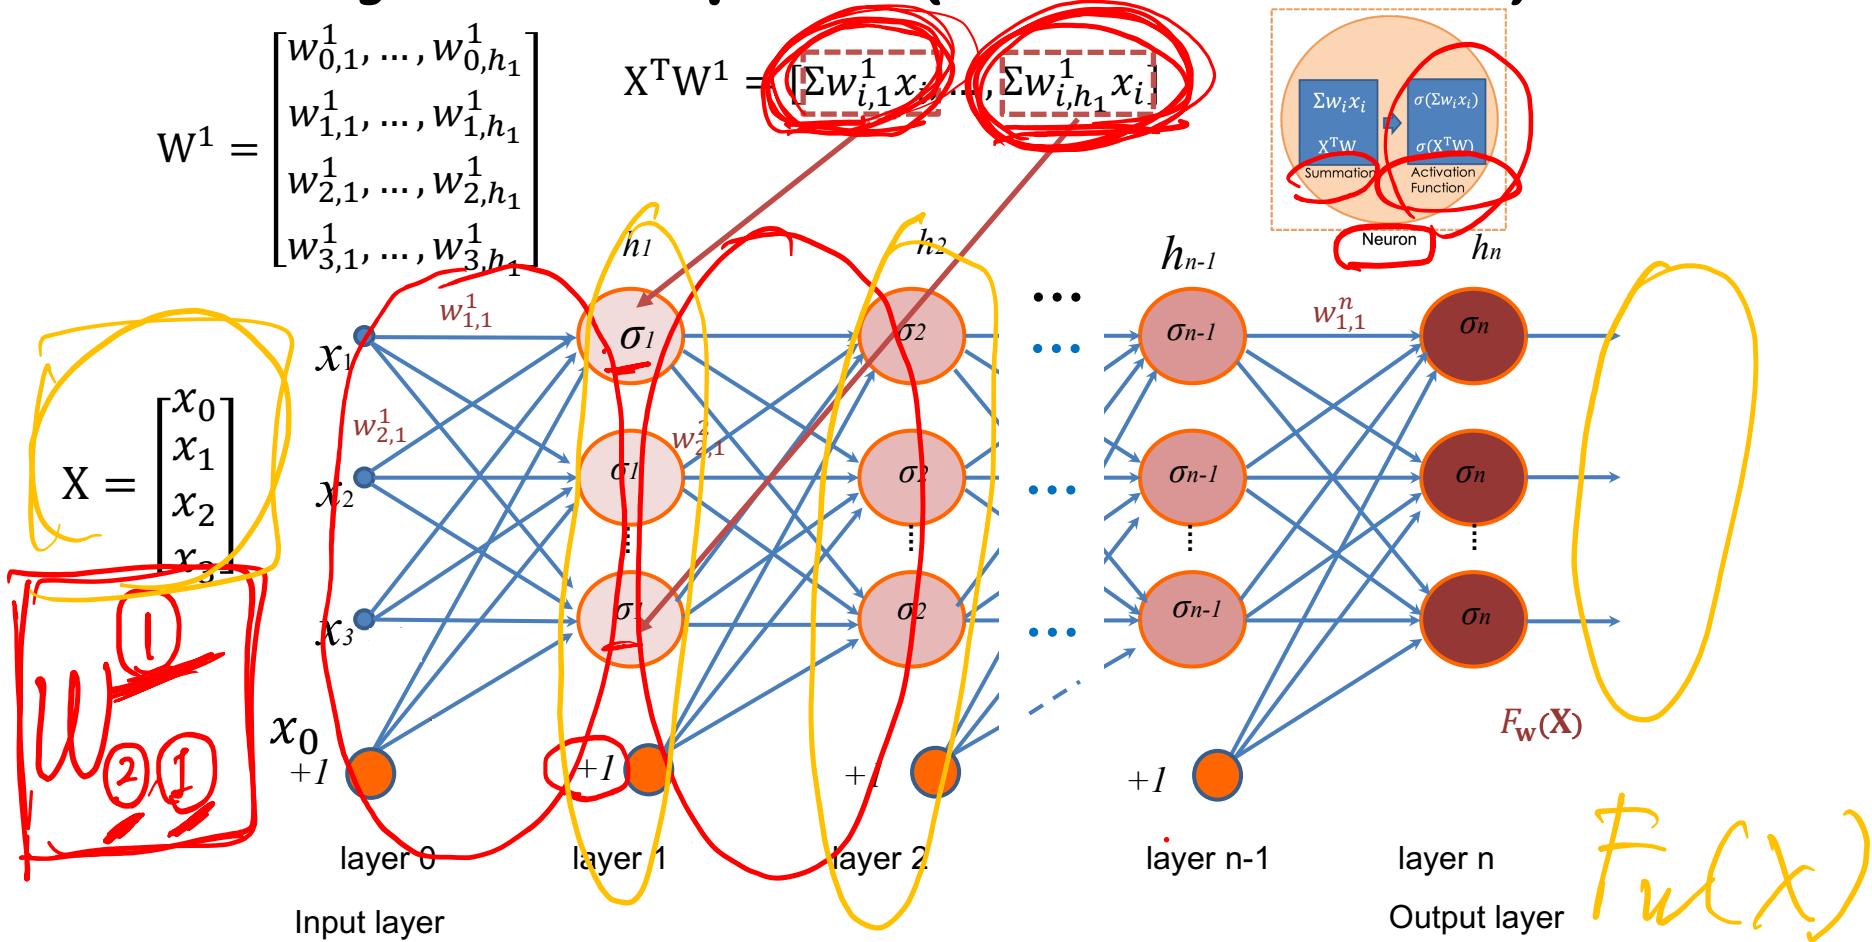


1. The number of hidden neurons in different layers may differ, i.e., h_1 don't have to be equal to h_2 .
2. For **classification** task, the number of neurons in the last layer equals to the number of classes.
3. We can treat the whole network as a function $F_w(X)$, where w is to be learned.

$F_w(X)$

parameters input

Multilayer Perceptron (Neural Network)



A neural network is essentially a **nested function**.

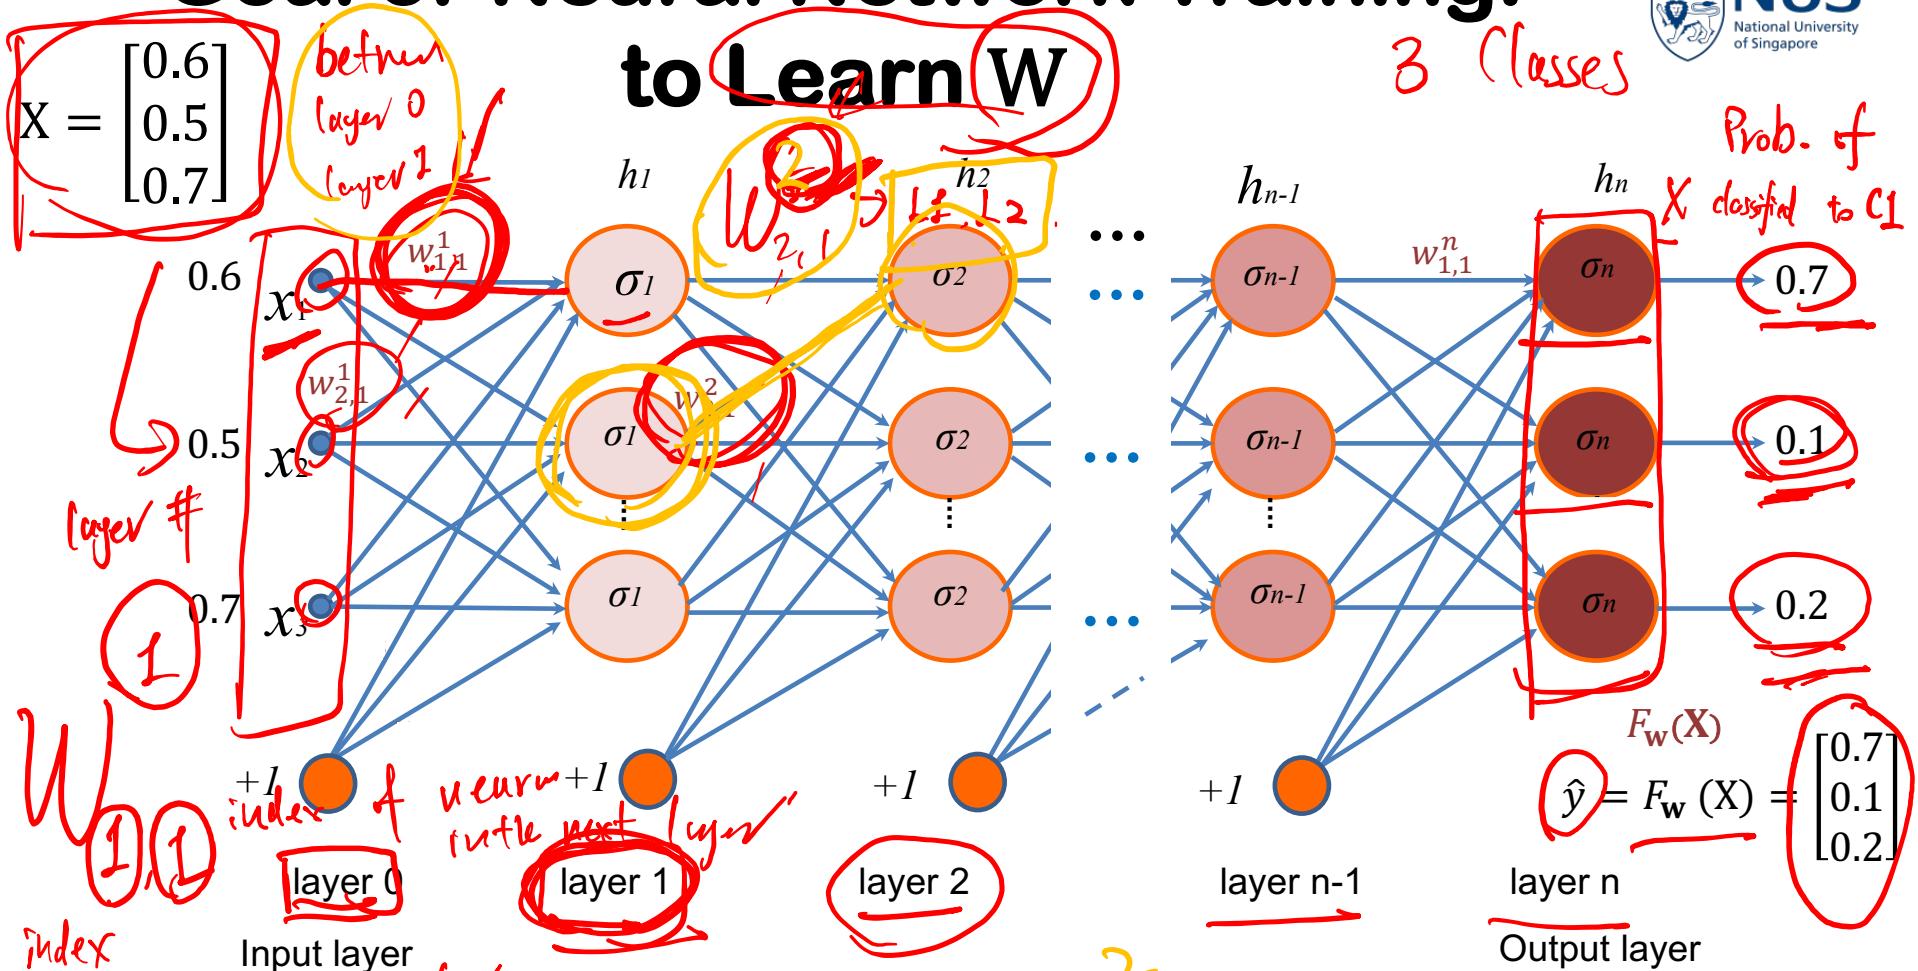
$$F_w(X) = \sigma([1, \dots, \sigma([1, \sigma(X^T W^1)] W^2) \dots] W^n)$$

Outline

- Introduction to Neural Networks
 - Multi-layer perceptron
 - Activation Functions
- Training and Testing of Neural Networks
 - Training: Forward and Backward
 - Testing: Forward
- Convolutional Neural Networks

Goal of Neural Network Training:

to Learn W

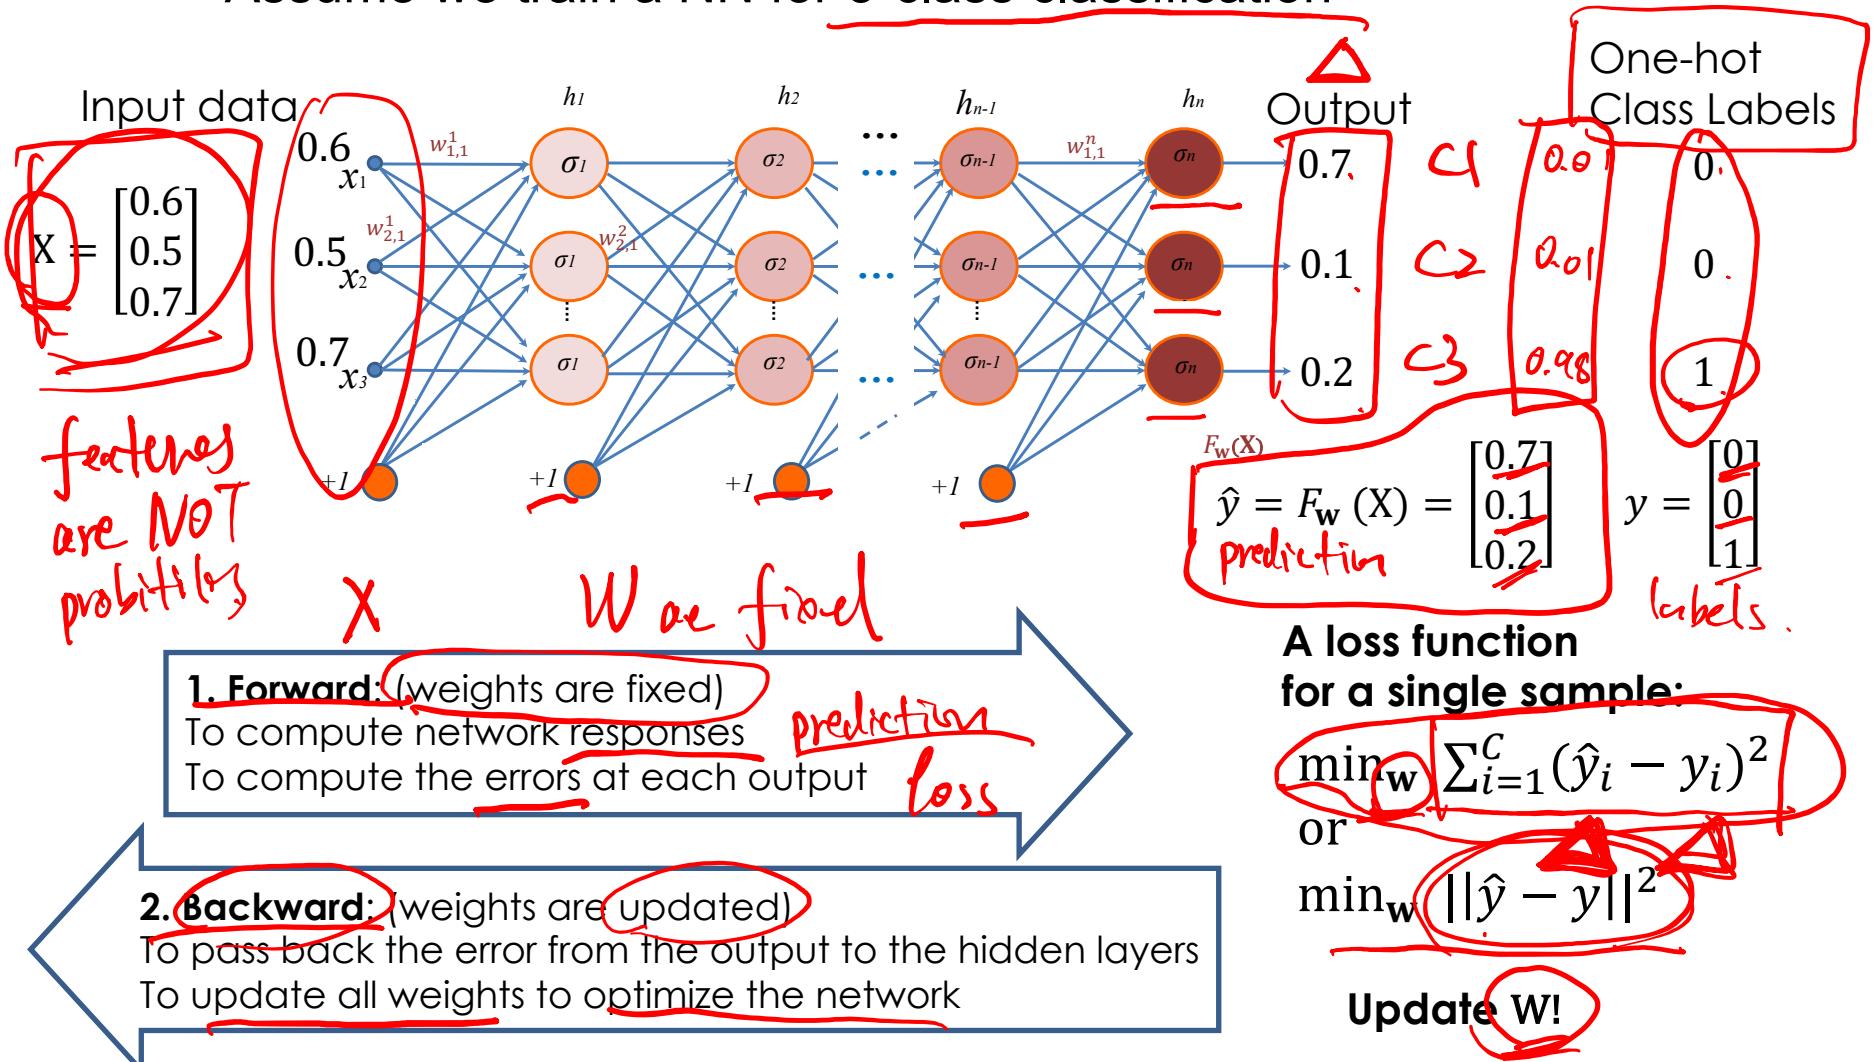


Specifically, W is learned through

1. Random initialization
2. Backpropagation

Neural Network Training: Backpropagation

Assume we train a NN for 3-class classification



Neural Network Training: Backpropagation

- Recall that the parameters W are randomly initialized.
- We use **Backpropagation** to update W .

- In essence, **Backpropagation** is gradient descent!

- Assume we have N samples, each sample denoted by X^j and the output of NN by \hat{y}^j , loss function is then

$$J = \sum_{j=1}^N \|\hat{y}^j - y^j\|^2, \quad \min_w J$$

of samples
 index samples

Overall loss function to W

Recall gradient descent in Lec 8: $w \leftarrow w - \eta \nabla_w J$

- We would therefore like to compute $\nabla_w J$!

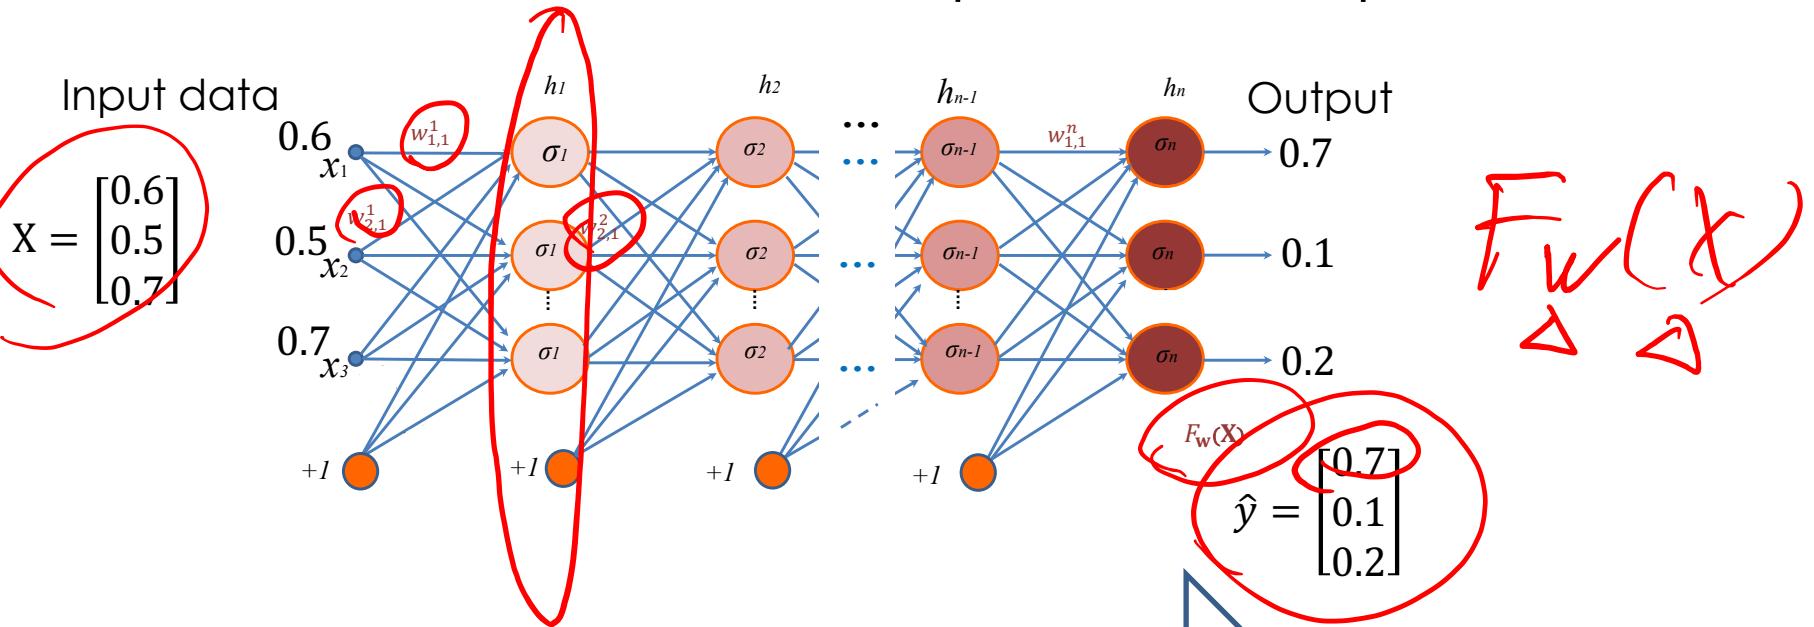
– J is a function of \hat{y} , and \hat{y} is a function of w , i.e., $\hat{y} = F_w(X)$

– Use gradient descent and chain rule!

Being aware of the concept is sufficient for exam. No calculation needed.

Neural Network Testing

Once all network is trained and parameters are updated

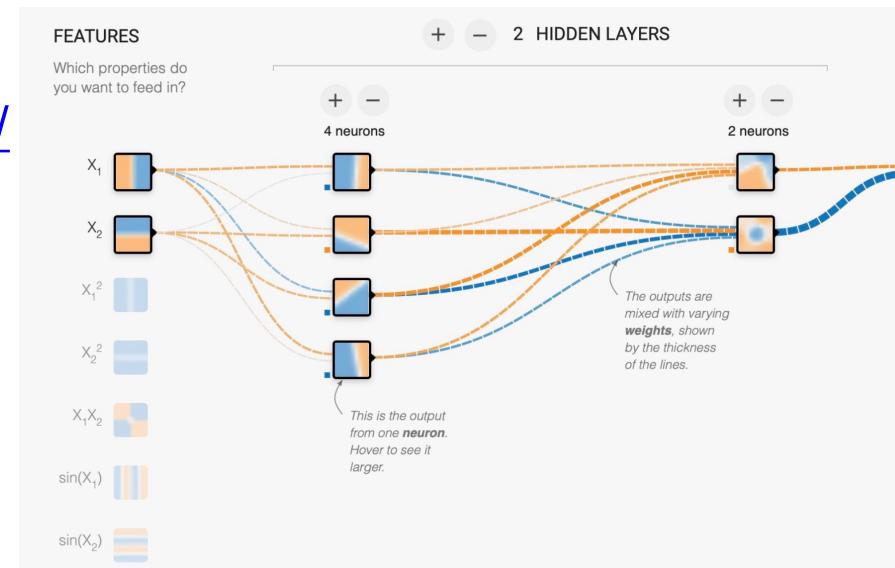


1. Forward: (weights are fixed)
 To estimate compute network responses
 To predict the output labels given novel inputs

Python Demo

```
[ ] from sklearn.neural_network import MLPClassifier
# Train the neural network classifier with 3 hidden layers of 200 neurons each
clf = MLPClassifier(hidden_layer_sizes=(200, 200, 200), activation="relu", learning_rate="invscaling", verbose=True)
clf.fit(X, y)
plot_clf(clf, X, y)
```

- Python Code (A simple multi-layer neural network classifier)
 - [lec12.ipynb](#)
- Demo
 - <https://playground.tensorflow.org/>



Supplementary materials (Not required for exam)

1) <https://www.youtube.com/watch?v=tIeHLnjs5U8>

This video series includes animations that explain backpropagation calculus.

2)

<https://www.youtube.com/playlist?list=PLQVvva0QuDcjD5BAw2DxE6OF2tius3V3>

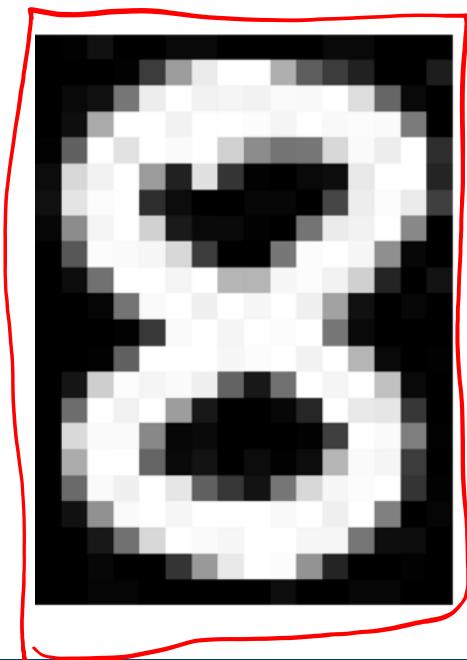
This video series includes hands-on coding examples in Python.

Outline

- Introduction to Neural Networks
 - Multi-layer perceptron
 - Activation Functions
- Training and Testing of Neural Networks
 - Training: Forward and Backward
 - Testing: Forward
- Convolutional Neural Networks

Convolutional Neural Network (CNN)

- A convolutional neural network (CNN) is a special type of neural network that significantly reduces the number of parameters in a deep neural network.
 - Very popular in image-related applications
 - Each image is stored as a matrix in a computer



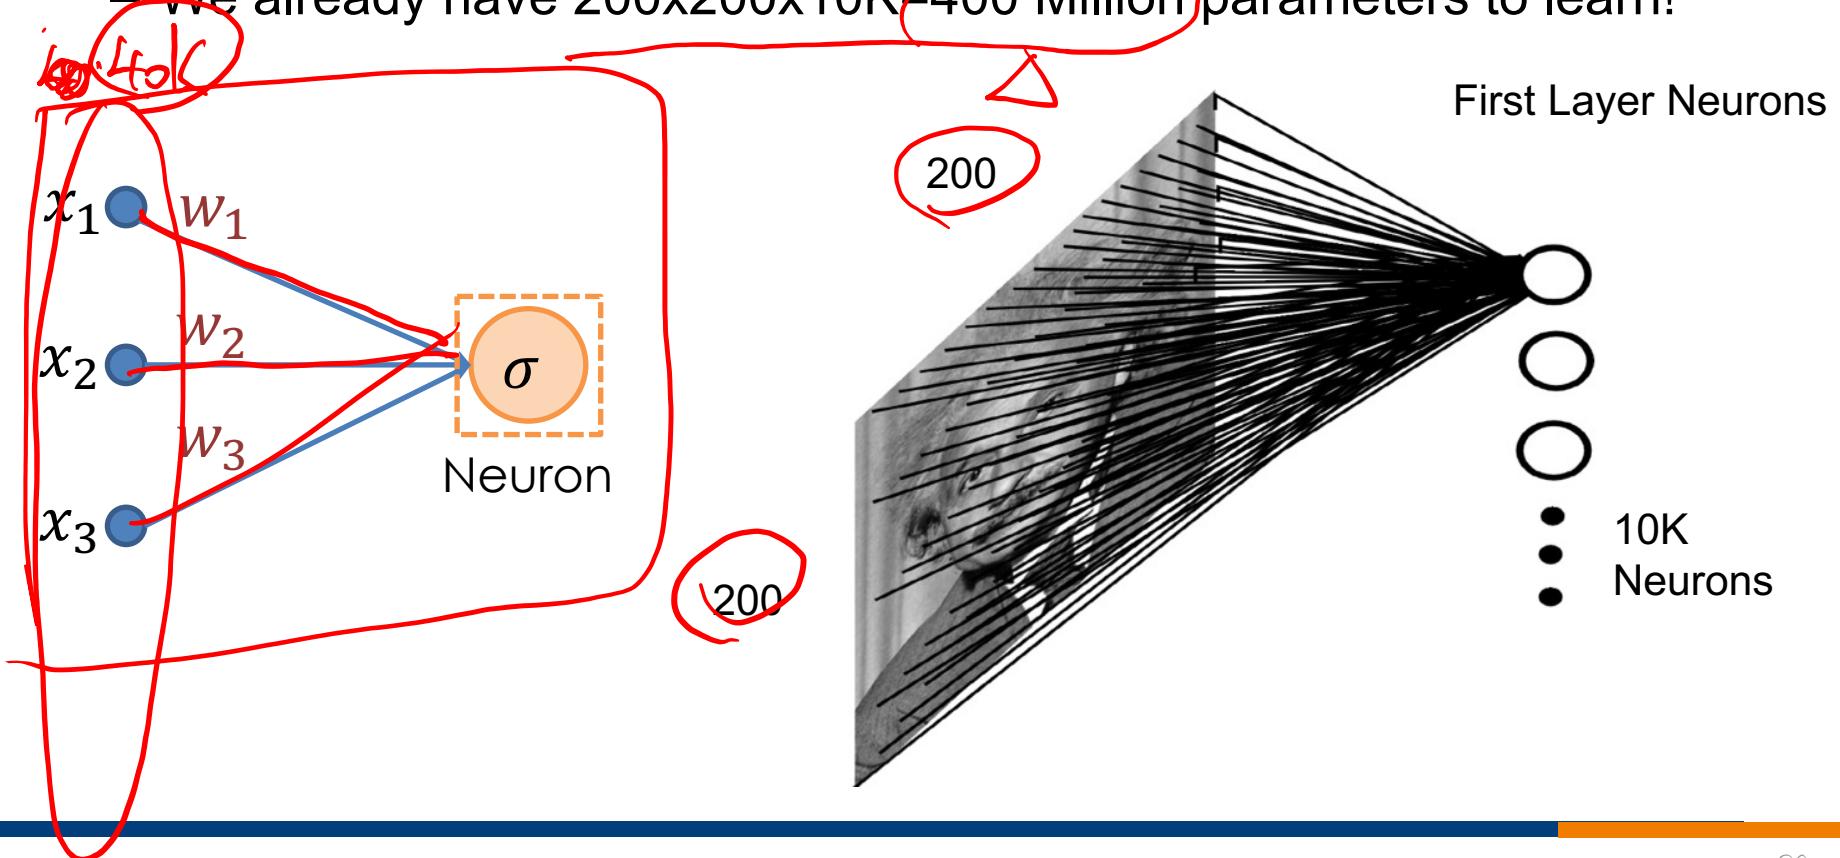
0	2	15	0	0	11	10	0	0	0	0	9	9	9	0	0	0
0	0	0	4	60	157	236	255	255	177	95	61	32	0	0	29	
0	10	1	11	238	255	244	245	243	250	249	255	222	103	10	0	
0	14	170	255	255	244	25	255	253	245	255	249	253	251	124	1	
2	98	255	228	255	251	254	211	141	116	122	215	251	238	255	49	
13	217	243	255	155	33	226	52	2	0	10	13	232	255	255	36	
16	229	252	254	49	12	0	0	7	7	0	70	237	252	235	62	
6	141	245	255	212	25	11	9	3	0	115	236	243	255	137	0	
0	87	252	250	248	215	60	0	1	121	252	255	248	144	6	0	
0	13	111	255	255	245	255	182	181	248	252	242	208	36	0	19	
1	0	5	117	251	255	241	255	247	255	241	162	17	0	7	0	
0	0	0	4	58	251	255	246	254	253	255	120	11	0	1	0	
0	0	4	97	255	255	255	248	252	255	244	255	182	10	0	4	
0	22	206	252	246	251	241	100	24	113	255	245	255	194	9	0	
0	111	255	242	255	158	24	0	0	6	39	255	232	230	56	0	
0	218	251	250	137	7	11	0	0	0	2	62	255	250	125	3	
0	173	255	255	101	9	20	0	13	3	13	182	251	245	61	0	
0	107	251	241	255	230	98	55	19	118	217	248	253	255	52	4	
0	18	146	250	255	247	255	255	255	249	255	240	255	129	0	5	
0	0	23	113	215	255	250	248	255	255	248	248	118	14	12	0	
0	0	6	1	0	52	153	233	255	252	147	37	0	0	4	1	
0	0	5	5	0	0	0	0	0	14	1	0	6	6	0	0	

0	2	15	0	0	11	10	0	0	0	0	0	9	9	9	0	0	0
0	0	0	4	60	157	236	255	255	177	95	61	32	0	0	0	29	
0	10	16	119	238	255	244	245	243	250	249	255	222	103	10	0	0	
0	14	170	255	255	244	25	25	253	245	255	249	253	251	124	1	0	
2	98	255	228	255	251	254	211	141	116	122	121	251	238	255	49	0	
13	217	243	255	155	33	226	52	2	0	10	13	232	255	255	36	0	
16	229	252	254	49	12	0	0	7	7	0	70	237	252	235	62	0	
6	141	245	255	212	25	11	9	3	0	115	236	243	255	137	0	0	
0	87	252	250	248	215	60	0	1	121	252	255	248	144	6	0	0	
0	13	113	255	255	245	255	182	181	248	252	242	208	36	0	19	0	
1	0	5	117	251	255	241	255	247	255	241	162	17	0	7	0	0	
0	0	0	4	58	251	255	246	254	254	253	255	120	11	0	1	0	
0	0	4	97	255	255	255	248	252	255	244	255	182	10	0	4	0	
0	22	206	252	246	251	241	100	24	113	255	245	255	194	9	0	0	
0	111	255	242	255	158	24	0	0	6	39	255	232	230	56	0	0	
0	218	251	250	137	7	11	0	0	0	2	62	255	250	125	3	0	
0	173	255	255	101	9	20	0	13	3	13	182	251	245	61	0	0	
0	107	251	241	255	230	98	55	19	118	217	248	253	255	52	4	0	
0	18	146	250	255	247	255	255	249	255	240	245	255	129	0	5	0	
0	0	23	113	215	255	250	248	255	255	248	248	118	14	12	0	0	
0	0	6	1	0	52	153	233	255	252	147	37	0	0	4	1	0	
0	0	5	5	0	0	0	0	0	14	1	0	6	6	0	0	0	

<https://medium.com/liteandtech/convert-csv-file-to-images-309b6fdb8c49>

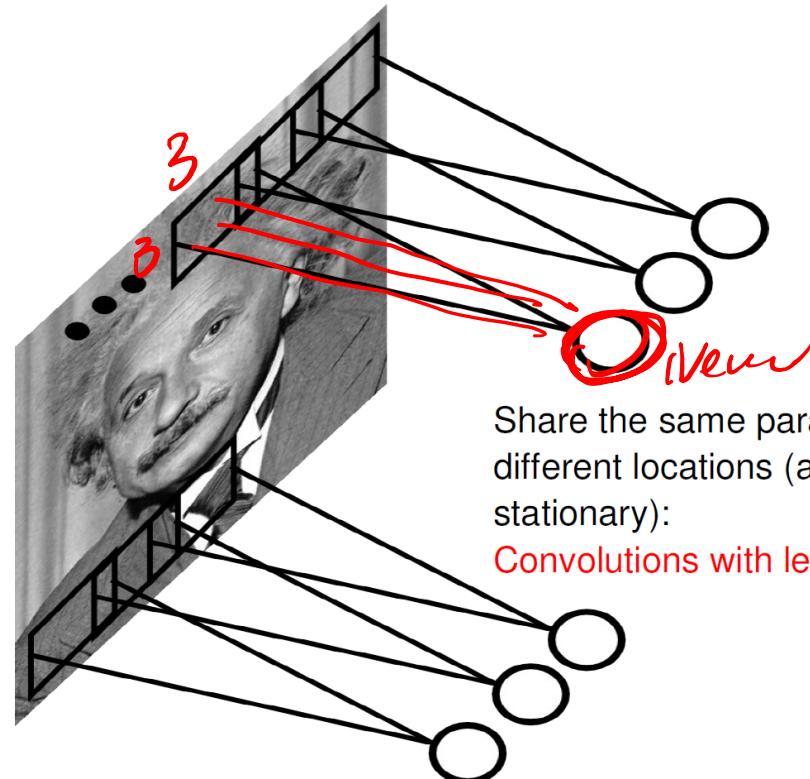
Convolutional Neural Network (CNN)

- If we model all matrix entries as inputs all at once
 - Assume we have an image/matrix size of 200×200 = $40K$
 - Assume we have $10K$ neurons in the first layer w_{ij}^k
 - We already have $200 \times 200 \times 10K = 400$ Million parameters to learn!



Convolutional Neural Network (CNN)

- Hence, we introduce CNN to reduce the number of parameters.
- Works in a sliding-window manner!



Convolutional Neural Network (CNN)

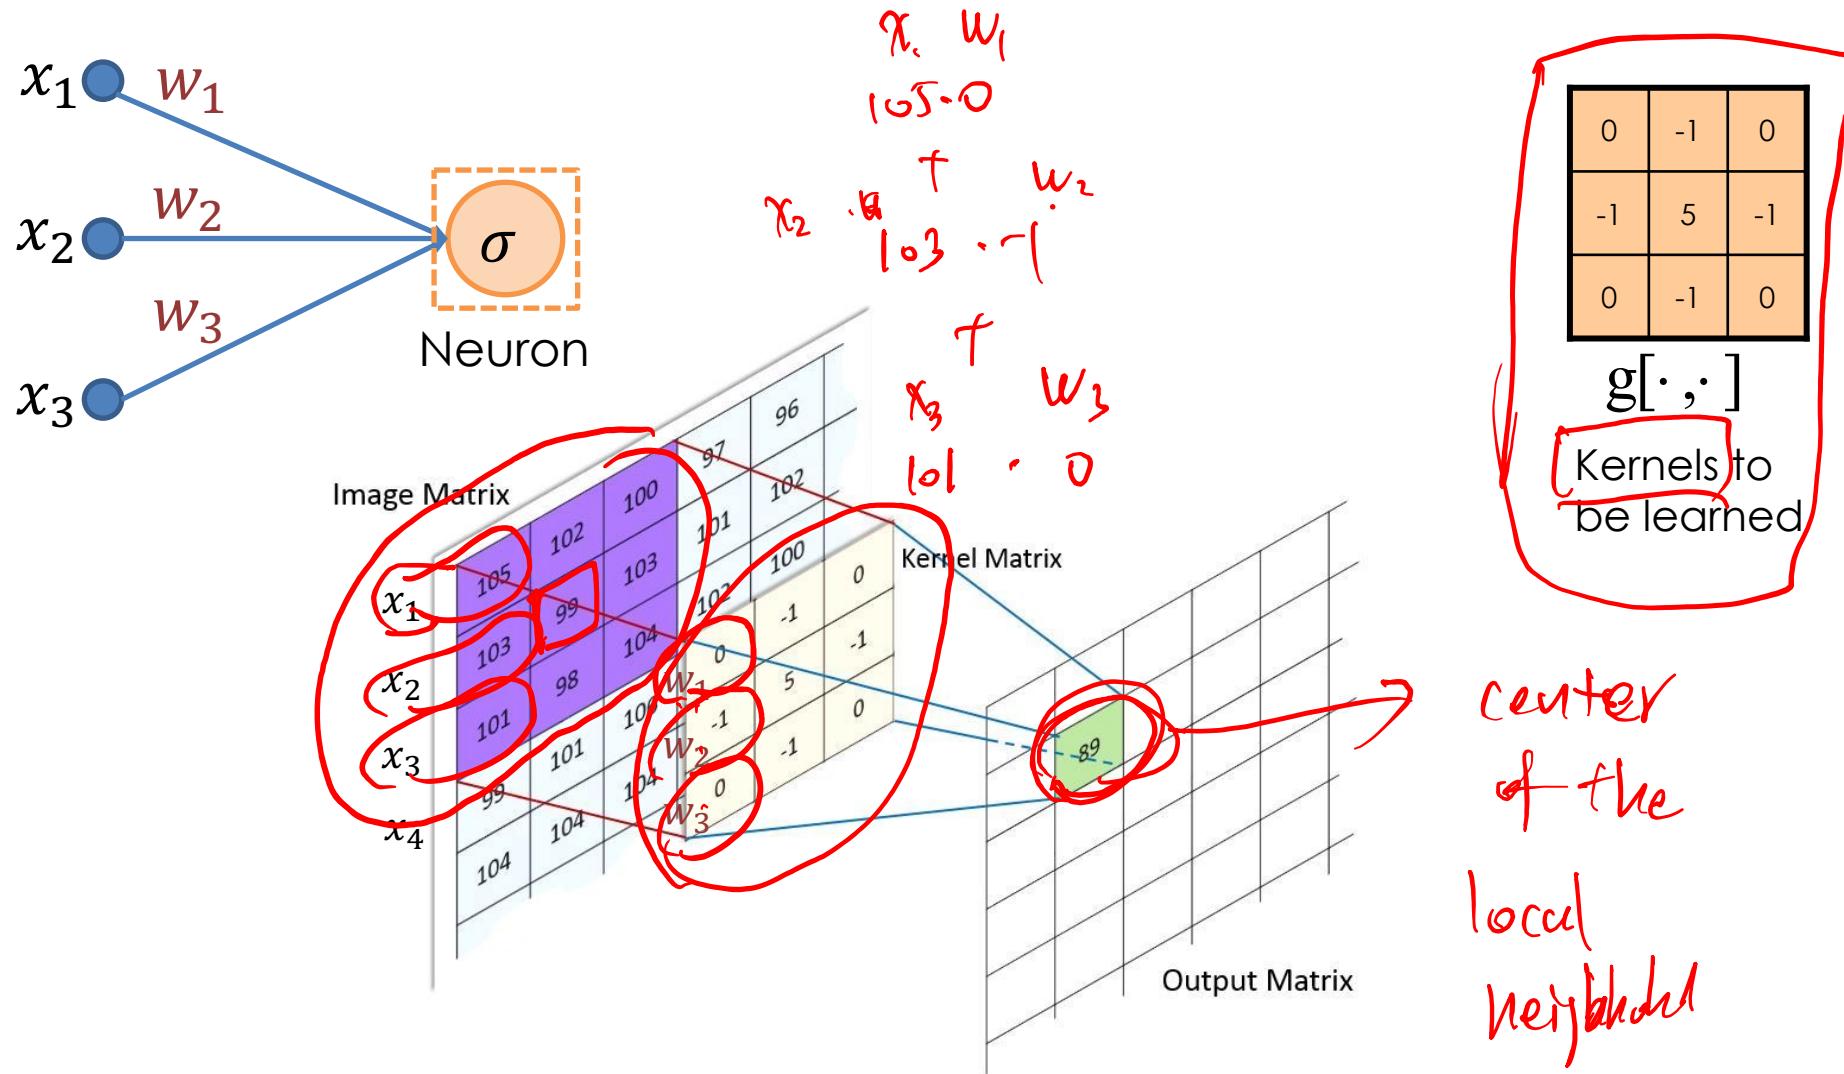
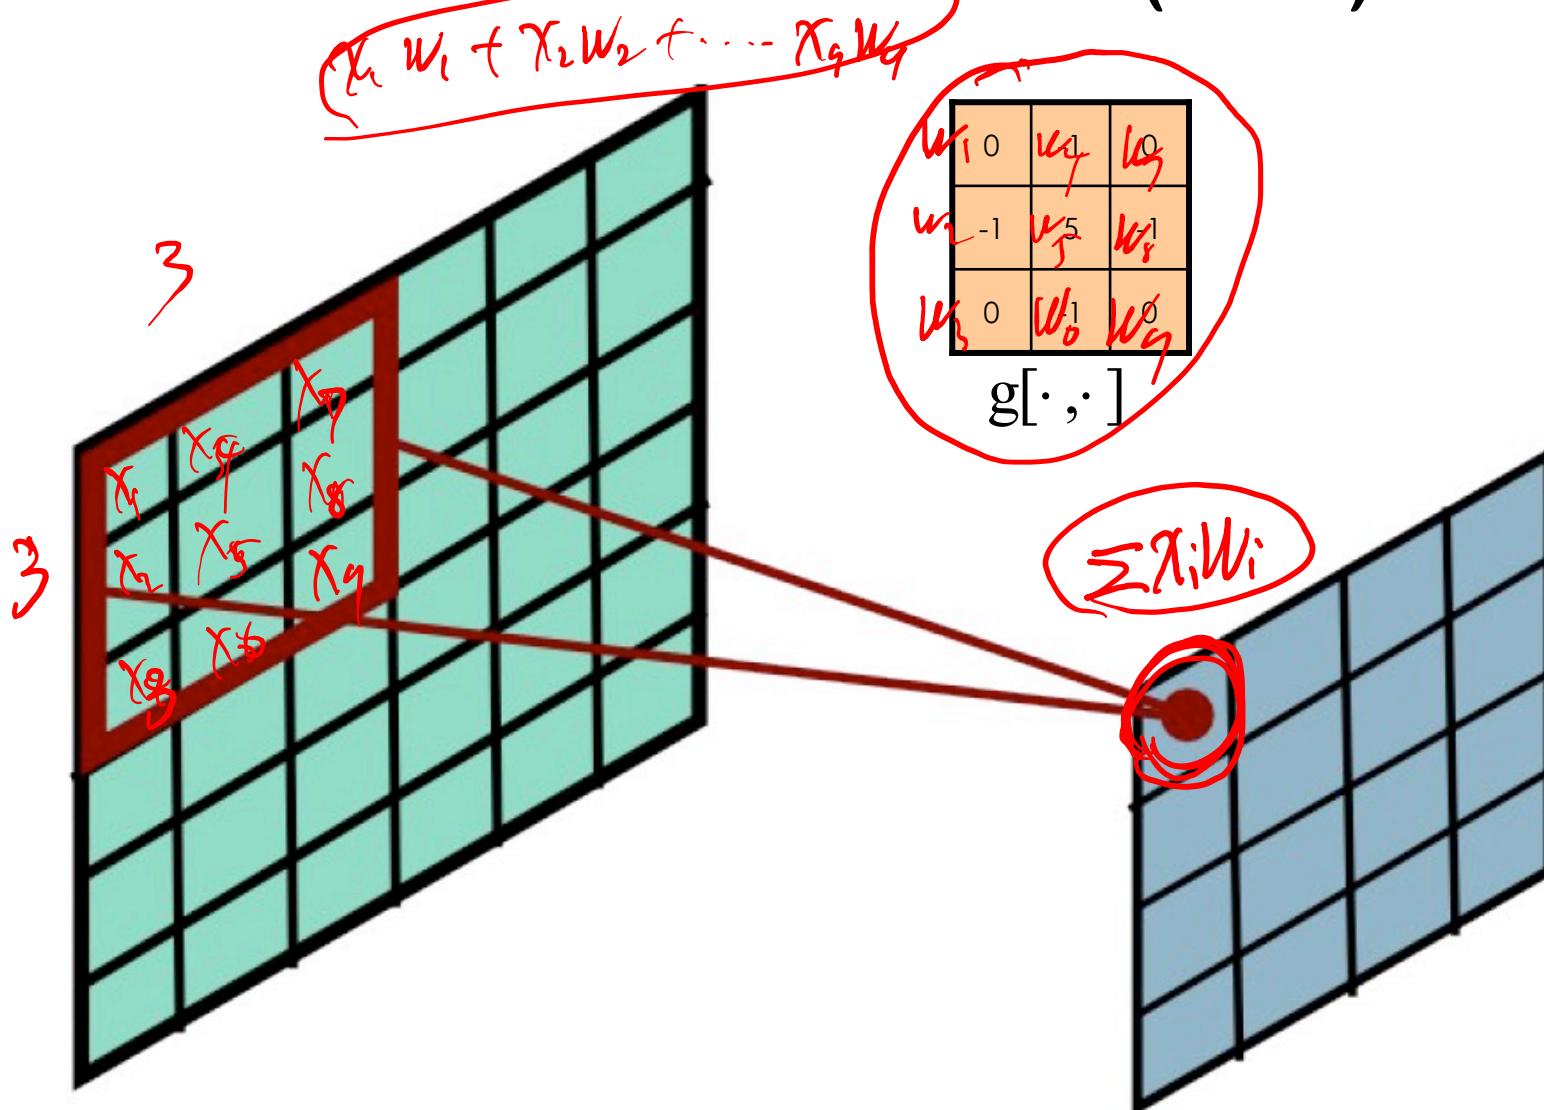
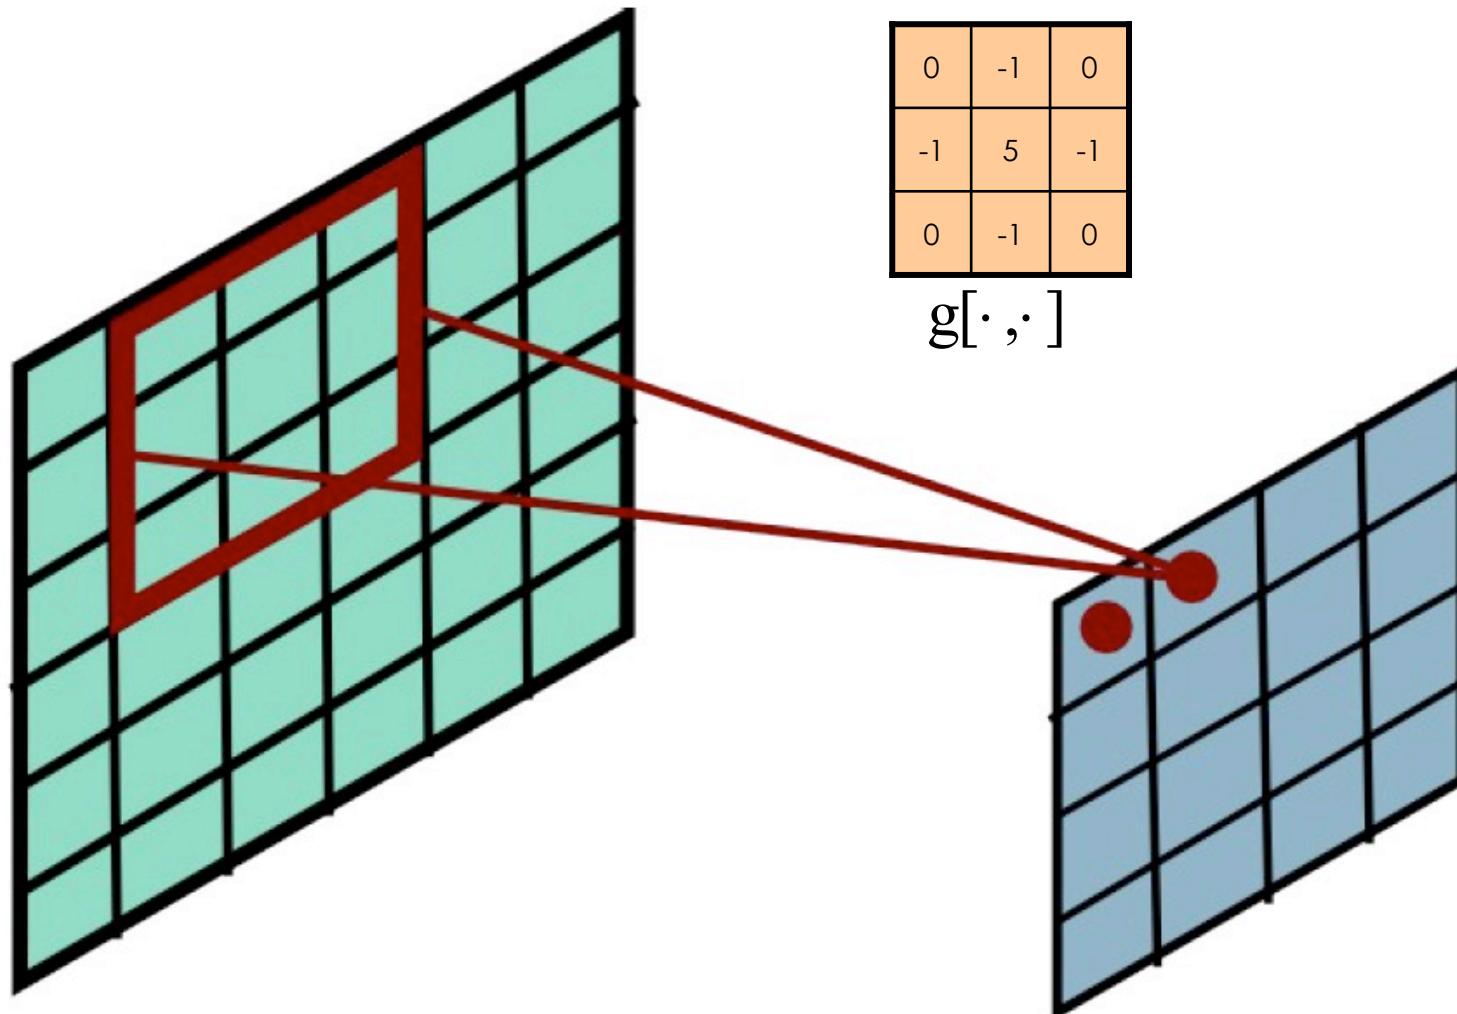


Image source: <https://brilliant.org/wiki/convolutional-neural-network/>

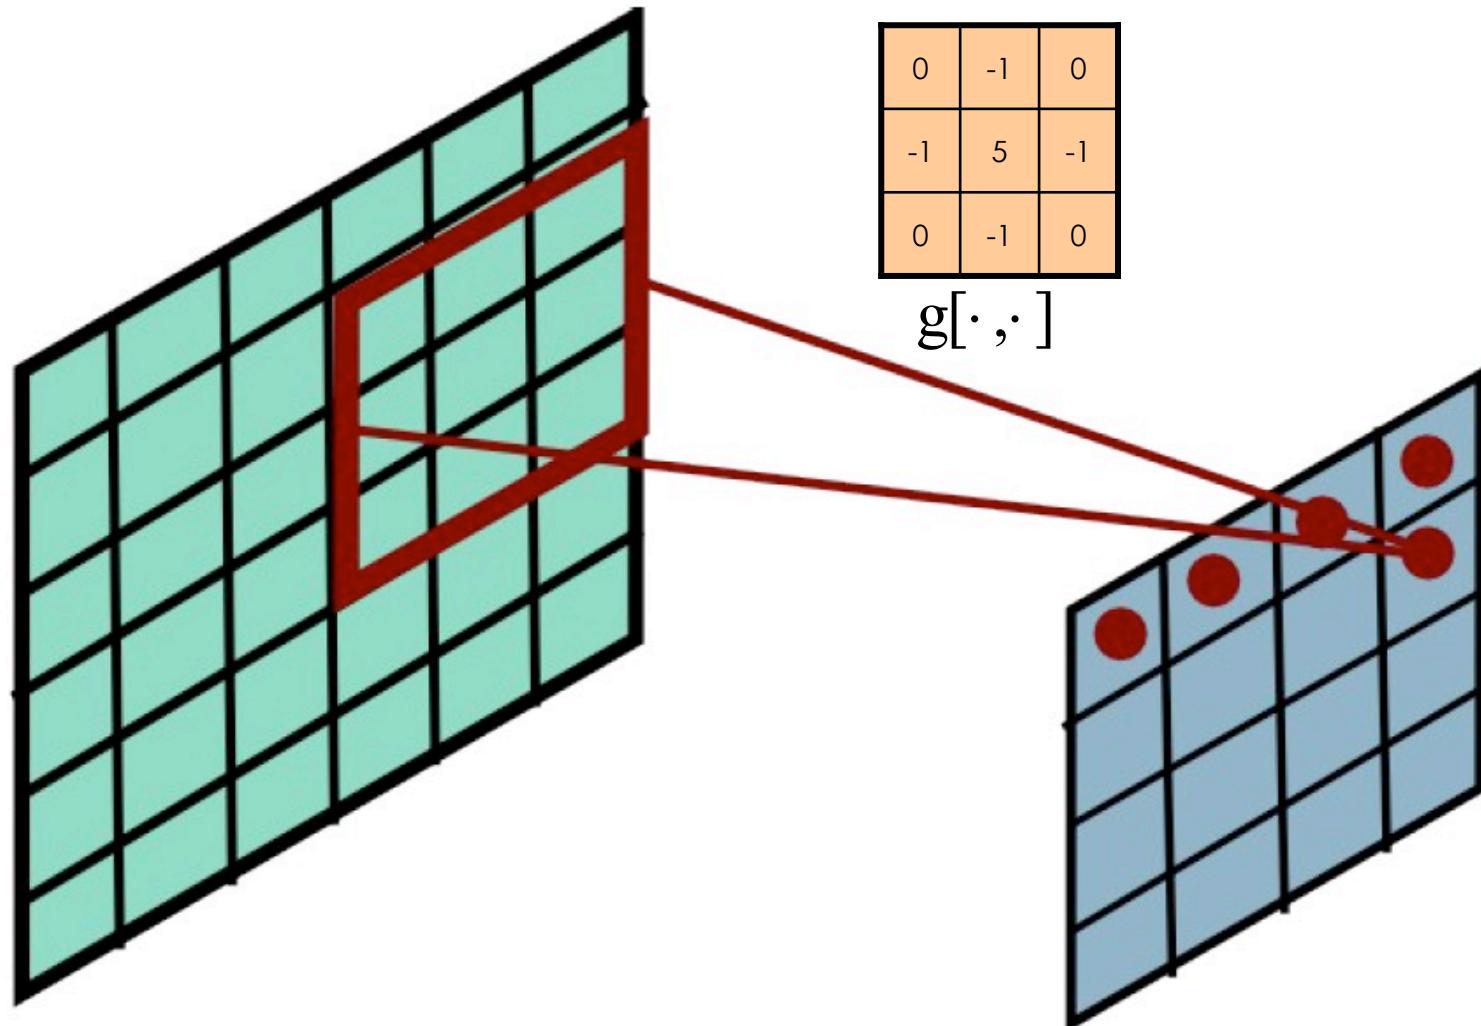
Convolutional Neural Network (CNN)



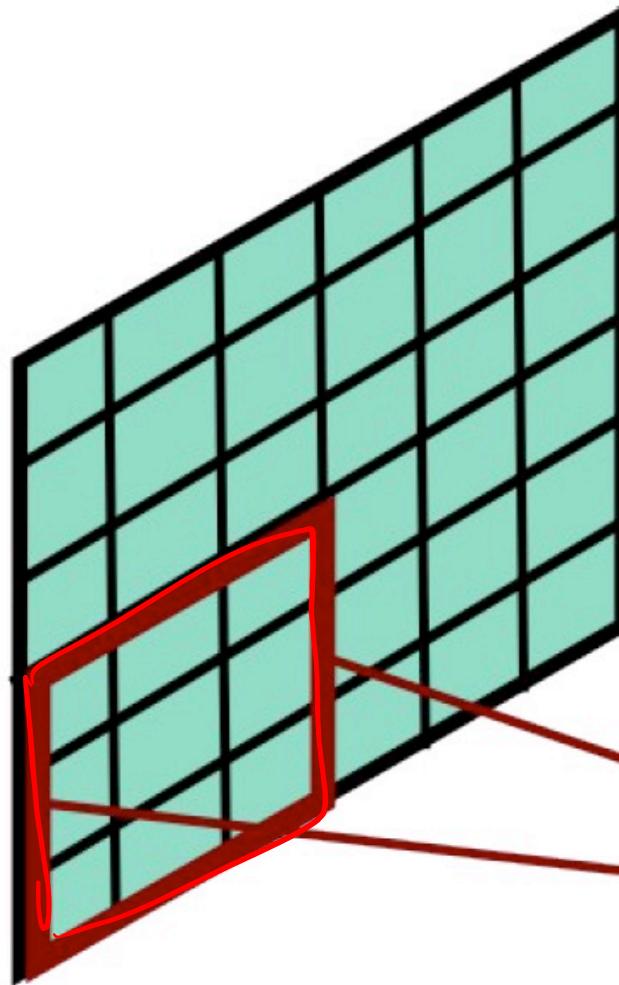
Convolutional Neural Network (CNN)



Convolutional Neural Network (CNN)

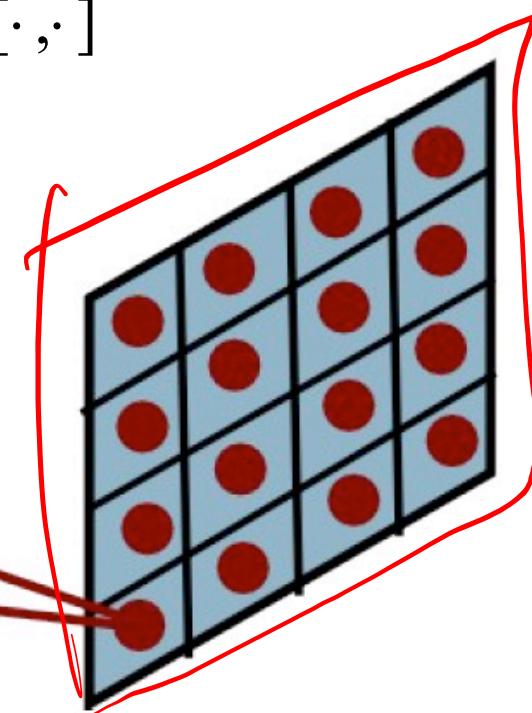


Convolutional Neural Network (CNN)



$$\begin{matrix} \text{W}_1 & & \\ \text{W}_2 & & \\ \text{W}_3 & & \end{matrix}$$
$$\begin{array}{|c|c|c|}\hline 0 & -1 & 0 \\ \hline -1 & 5 & -1 \\ \hline 0 & -1 & 0 \\ \hline \end{array}$$

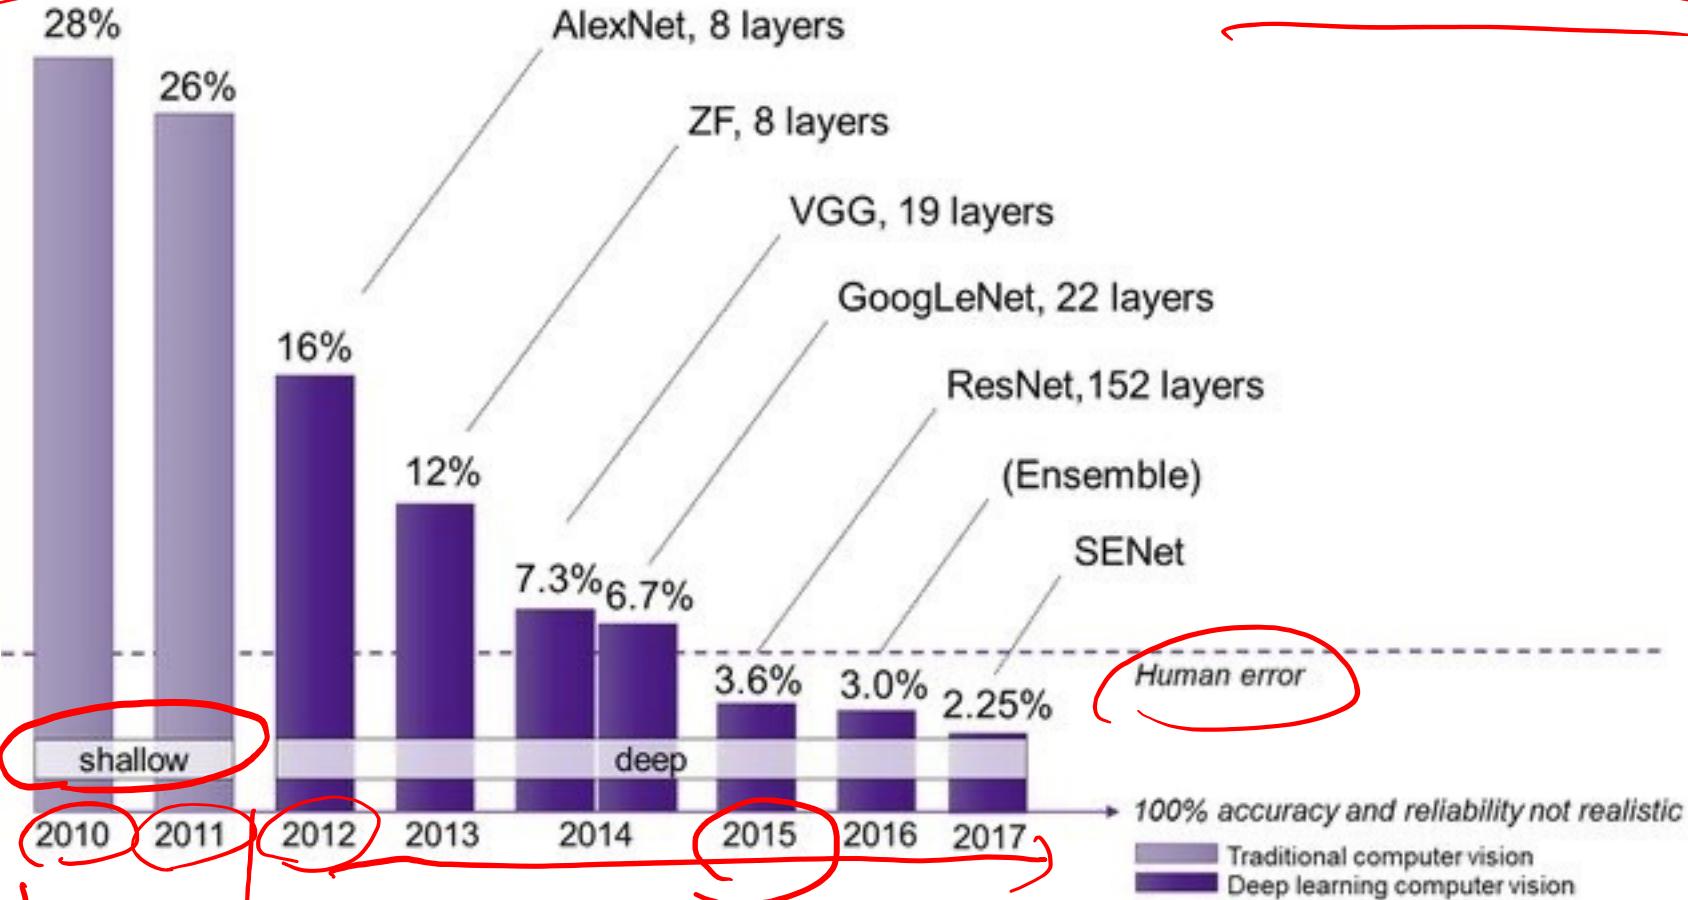
$g[\cdot, \cdot]$



Neural Networks are Effective

Imagret

Error Rate



not
neural networks

Summary

- Introduction to Neural Networks
 - Multi-layer perceptron
 - Activation Functions
- Training and Testing of Neural Networks
 - Training: Forward and Backward
 - Testing: Forward
- Convolutional Neural Networks

(CNN)

Thanks everyone, for your time and effort throughout the semester!

